



Department of Mathematics and Computer Science
Coding Theory and Cryptology Group

Secure Sessions for Ad Hoc Multiparty Computation in MPyC

Master's thesis

Emil Nikolov

Id nr: 0972305
`emil.e.nikolov@gmail.com`

Supervisor : Dr. ir. L.A.M. (Berry) Schoenmakers

June 13, 2023

Contents

Contents	iii
Todo list	iv
Glossary	vii
List of Figures	ix
1 Introduction	1
I The State of Multiparty Communications over the Internet	3
2 Internet Communications	7
2.1 Communication Protocols	8
2.1.1 Network Layer (L3) Communication	8
2.1.2 Transport Layer (L4) Communication	8
2.1.3 Application Layer (L7) Communication	9
2.2 Secure Communication Protocols	9
2.2.1 Transport Layer (L4) to Application Layer (L7) Security	9
2.2.2 Network Layer (L3) Security Protocols	10
2.3 Multiparty Communication - Barriers and Traversal	12
3 Overlay Networks	17
3.1 Network layer (L3) and Data Link (L2) Overlay Networks	19
3.1.1 Traditional VPNs	19
3.1.2 Mesh VPNs	20
3.2 Application Layer (L7) Network Overlays	20
3.2.1 OpenZiti	20
II Proof of Concepts and Evaluation	23
4 Testing methodology	25
4.1 Measuring performance	25
4.2 Security	26
4.3 Usability	26

5	Internet Protocol based solution	27
5.1	Implementation details	27
5.2	Performance analysis	27
5.3	Security analysis	27
5.4	Usability analysis	27
6	WireGuard based solution	29
6.1	Implementation details	29
6.2	Performance analysis	29
6.3	Security analysis	29
6.4	Usability analysis	29
7	Tailscale based solution	31
7.1	Implementation details	31
7.2	Performance analysis	31
7.3	Security analysis	31
7.3.1	Trust model	31
7.3.2	Identity	31
7.4	Usability analysis	31
8	Headscale based solution	33
8.1	Implementation details	33
8.2	Performance analysis	33
8.3	Security analysis	33
8.3.1	Trust model	33
8.3.2	Identity	33
8.4	Usability analysis	33

Todo list

include a section on MPC?	5
Is there a more fitting section name than "The Internet"	7
add a paragraph that describes what the rest of the section will contain	8
packet is the physical envelope of the IP datagram	8
show a diagram of the internet with multiple local networks	8
VPN This way the VPN client can hide the IP addresses of the final destinations of the traffic from hosts between itself and the VPN server. Depending on the use case, the VPN server could either facilitate the client's secure access to other resources on the virtual network, or provide anonymity to the VPN client by hiding its IP address from the hosts its communicating with on the public Internet.	11
The Internet Protocol Security (IPSec) protocol suite is designed specifically to provide secure communication between two hosts at the network layer (L3) of the OSI model. IPSec can be used to secure traffic between two hosts, between a host and a network, or between two networks. It provides a range of services, including data confidentiality, data integrity, and authentication, and is often used in Virtual Private Network (VPN) connections to create secure tunnels between remote networks. IPSec operates in either Transport or Tunnel mode, and can use either the Authentication Header (AH) or Encapsulating Security Payload (ESP) protocols to provide its security services.	11
- Internet Protocol Security (IPSec) - Layer 3 protocol suite part of the Internet Protocol Suite - used inside VPN software - has implementations in both user and kernel space as well as hardware implementations - rewrites and encrypts the IP headers and payloads - virtual routing table - Initially was built into IPv6, separate from IPv4	11
elaborate	11
noise is transport agnostic	11
noise has limited cipher suites	11
Classless Inter-Domain Routing (CIDR) is a common notation for describing IP address ranges, e.g. '192.168.0.1/16', where the number after the slash describes the bit-length of the fixed prefix for a subnet.	12
describe some of the nat behaviors, e.g. if the source IP address/port vary per destination are changed depending on the destination/port mapping algorithms, if it maps ports, IPs, whether the mapped IPs are different per destination and others.	13
connection reversal	14
If, else?	14
Talk about traversal friendly NATs and unfriendly NATs	14

CONTENTS

■ add a bullet list with the steps in the STUN process and relate them to the step numbers in the figure	14
■ only 65000 ports per IP address means that CGNATs that provide more than 65000 connections from client devices require more than one public IP address	15
■ hairpinning - Hairpinning, also known as NAT loopback or NAT reflection, is a technique used by NAT devices to allow hosts on a private network to access a public server using its public IP address. Without hairpinning, the NAT device would not recognize the connection as a loopback connection and would route it to the public network, causing the connection to fail. With hairpinning, the NAT device recognizes that the connection is a loopback connection and redirects the traffic back to the same NAT device, which then forwards the traffic to the correct host on the private network. This can be useful in scenarios where a private network is hosting a public-facing server that is also accessed by internal users on the same network using its public IP address.	15
■ webrtc	15
■ explain the names	19
■ tls vs ipsec vpns. TLS vpns offer a virtual network interface at layer 3, but run over L7 TLS	19

Glossary

***E*³** Extensible Evaluation Environment [25](#)

ACL Access Control List [20](#)

CA Certificate Authority [10](#)

CGNAT Carrier-Grade NAT [14](#)

CIDR Classless Inter-Domain Routing [v](#), [12](#)

DERP Designated Encrypted Relay for Packets [13](#), [20](#)

DNS Domain Name System [9](#)

DTLS Datagram Transport Layer Security [9](#)

ECDH Elliptic Curve Diffie-Hellman [11](#)

FTP File Transfer Protocol [8](#)

HTML HyperText Markup Language [8](#)

HTTP HyperText Transfer Protocol [9](#), [10](#)

HTTPS HyperText Transfer Protocol Secure [10](#)

ICE Interactive Connectivity Establishment [15](#)

IP Internet Protocol [8](#), [11](#), [19](#)

IPSec Internet Protocol Security [v](#), [11](#)

ISP Internet Service Provider [7](#), [10](#), [14](#)

LAN Local Area Network [8](#), [19](#)

NAT Network Address Translation [12](#)

NAT-PMP NAT Port Mapping Protocol [13](#)

- OSI** Open Systems Interconnection [5](#)
- P2P** Peer to Peer [19](#), [20](#)
- PCP** Port Control Protocol [13](#)
- PDU** Protocol Data Unit [8](#)
- PKI** Public Key Infrastructure [10](#)
- PSK** Pre-Shared Key [11](#)
- QUIC** Quick UDP Internet Connections [9](#), [10](#)
- SMTP** Simple Mail Transfer Protocol [8](#)
- SSH** Secure Shell Protocol [8](#), [10](#), [11](#)
- SSL** Secure Sockets Layer [9](#)
- STUN** Session Traversal Utilities for NAT [14](#)
- TAP** Network Tap [19](#)
- TCP** Transmission Control Protocol [8](#), [9](#), [10](#)
- TLS** Transport Layer Security [9](#), [11](#), [17](#)
- TUN** Network Tunnel [19](#)
- TURN** Traversal Using Relays around NAT [13](#)
- UDP** User Datagram Protocol [8](#), [9](#), [14](#), [17](#)
- UPnP** Universal Plug and Play [13](#)
- URL** Universal Resource Locator [9](#)
- VPN** Virtual Private Network [v](#), [10](#), [11](#), [17](#), [19](#), [20](#)
- WebRTC** Web Real-Time Communication [15](#)
- WWW** World Wide Web [8](#)

List of Figures

2.1	OSI model mapping of the Internet Protocol Suite	7
2.2	Two parties behind separate NATs	13
2.3	NAT traversal via STUN	14
3.1	OSI model mapping of various protocols	18

LIST OF FIGURES

Chapter 1

Introduction

Part I

The State of Multiparty Communications over the Internet

This chapter provides background information on the challenges of Internet communications between multiple independent parties. It also presents a systematic overview of the available solutions, using the [Open Systems Interconnection \(OSI\)](#) reference model as a conceptual framework. Section [2](#) briefly explores the fundamentals of the Internet, its protocols, the limitations for peer-to-peer protocols, and some of the approaches to overcome them. Section [3](#) discusses higher-level overlay networks that build on top of the lower-level protocols from section [2](#).

include a
section on
MPC?

Chapter 2

Internet Communications

The Internet is a global network that consists of numerous interconnected computer networks spanning billions of host devices owned by diverse parties from around the world. Key components of the Internet include the Internet Protocol Suite (known as TCP/IP) and the physical infrastructure that connects the individual networks. Sections of the infrastructure are deployed and managed by different tiers of [Internet Service Providers \(ISPs\)](#) who also maintain links between each other. To ensure efficient utilization of the hardware, the Internet relies on packet-switching techniques that divide the data traffic into smaller individually processed packets.

Is there a more fitting section name than "The Internet"

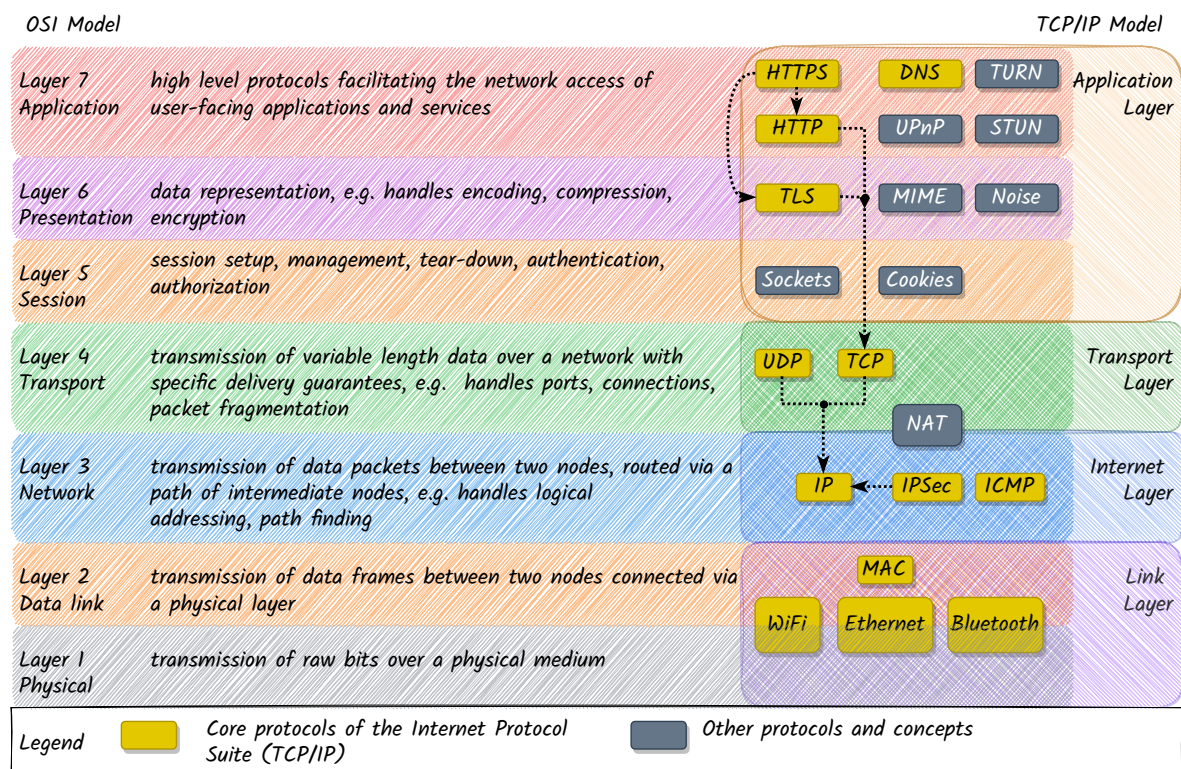


Figure 2.1: OSI model mapping of the Internet Protocol Suite

Communication protocols are usually organized into abstraction layers based on the scope of their functionality. Several reference models define different layering schemes, for example, the OSI model recognizes 7 layers, while TCP/IP itself combines some of the layers and recognizes 4. Figure 2.1 shows how the two models relate to each other and describes the responsibilities of the various layers. Throughout this thesis, we will refer to the 7 layer numbers of the OSI model as they are more widely used in the literature.

Each communication protocol has its own **Protocol Data Unit (PDU)** that

Services that are implemented as Application layer (L7) protocols on top of TCP/IP include the **World Wide Web (WWW)**, file transfer (**FTP**), email (**SMTP**), instant messaging, remote access (**SSH** [LY06]) and others. The Web is a collection of interconnected documents that use Web technologies such as **HTML** and JavaScript. It is typically accessed via a user-agent software such as a **Web Browser**.

The following sub-sections will briefly cover the main protocols of the Internet Protocol Suite, the issues for multiparty communications and some of the low-level mitigation techniques.

2.1 Communication Protocols

2.1.1 Network Layer (L3) Communication

The **IP** [81] is a Network layer (L3) protocol of the Internet Protocol Suite that is responsible for transferring datagrams between devices across the boundaries of their **Local Area Networks (LANs)** by possibly routing them via multiple intermediate devices (e.g. routers). A datagram is a self-contained unit of data, typically associated with connectionless protocols that provide no guarantees for delivery or ordering (e.g. IP, UDP). IP datagrams have a header that contains fields such as the **IP addresses** of its source and destination, and a payload that encapsulates the data from the Transport Layer (L4) protocols. A **router** is a device that is part of multiple networks and relays datagrams between them based on a routing table that maps IP address ranges to networks.

2.1.2 Transport Layer (L4) Communication

User Datagram Protocol (UDP) is a very thin Transport layer (L4) protocol that only provides port multiplexing and checksumming on top of IP. - Port multiplexing - uses 16-bit numbers to allow multiple processes behind the same IP address to establish their own communication channels - Checksumming - used to detect errors in the datagram header and payload

As with IP, UDP packets are referred to as datagrams because they are not delivered reliably and if such features are required, they must be implemented by the consumer of the protocol.

Transmission Control Protocol (TCP) is another Transport layer (L4) protocol. Like UDP, it provides port multiplexing and checksumming, but it offers stronger delivery guarantees. Some of the features it offers are listed below:

- Connection management - TCP establishes reliable connections between the communicating hosts and can gracefully terminate them when required
- Segmentation - TCP splits variable-length data streams into segments that fit inside IP datagrams and transmits them individually

add a paragraph that describes what the rest of the section will contain

packet is the physical envelope of the IP datagram

show a diagram of the internet with multiple local networks

- ordering - segments have sequence numbers to ensure that they are reassembled in the correct order at the receiving host
- Error detection and correction - TCP retransmits a segment if its checksum fails

Both TCP and UDP are useful in different scenarios. UDP is faster and is used for applications that can tolerate packet loss, e.g. video streaming, VoIP, or in cases where it is preferable for an application to implement its own reliable delivery. TCP has a higher overhead than UDP but its reliable delivery is a good default for most applications on the Internet.

[Quick UDP Internet Connections \(QUIC\)](#) is a more recent Transport layer (L4) protocol that is built on top of UDP and offers similar features to TCP, but with lower latency in some scenarios. It provides an additional level of multiplexing within a process, which allows multiple streams of data to be sent over the same connection asynchronously. To better optimize the establishment of secure connections, QUIC is tightly coupled with [TLS](#) version 1.3. One of the main uses of QUIC is inside [HTTP](#) version 3.

2.1.3 Application Layer (L7) Communication

[HyperText Transfer Protocol \(HTTP\)](#) is an Application layer (L7) protocol that enables stateless request/response interactions on the Web between web servers and clients (e.g. browsers). Similar to other L7 protocols, it uses [Universal Resource Locators \(URLs\)](#) for locating resources.

URL format: `scheme://host:port/path?query=value#fragment`

URL example: `http://www.example.com:80/path/to/file.html`.

HTTP up to version 2 uses TCP as a transport protocol and since version 3 uses [QUIC](#). HTTP provides several features such as:

- Request Methods - used by the client to specify the action to perform on the resource behind the given URL, e.g. GET, POST, PUT, DELETE, etc.
- Headers - used to provide additional information about a request or response, e.g. Content-Type, Authorization, Cache-Control
- Status codes - used to indicate the result of a request, e.g. if it was successful (200), or if the resource is missing (404)
- Cookies - used to include stateful information about the user kept on the client-side
- Caching - used to specify that the result of a request can be cached for a certain time to avoid repeating the request's action.

The [Domain Name System \(DNS\)](#) operates at the Application Layer (L7) and allows the conversion of human-readable domains to IP addresses, e.g. `google.com` to `142.250.179.142`.

2.2 Secure Communication Protocols

2.2.1 Transport Layer (L4) to Application Layer (L7) Security

[Transport Layer Security \(TLS\)](#) [Res18] and its precursor [Secure Sockets Layer \(SSL\)](#) are protocols that provide secure communications to Application layer (L7) protocols on top of a reliable Transport layer (L4) protocol like [TCP](#). [Datagram Transport Layer Security \(DTLS\)](#) is a related protocol that works with connectionless transport protocols like [UDP](#). TLS is

usually placed somewhere between the Presentation layer (L6) and the Transport layer (L4) because Application layer (L7) protocols use it as a transport protocol while having to manage the TLS connections.

TLS relies on digital certificates and [Public Key Infrastructure \(PKI\)](#) to establish trust between the communicating parties and to prevent man-in-the-middle attacks. A certificate includes information such as:

- Subject - an identifiable name for the certificate's owner. Depending on the use case it can be a domain name, an IP address, an email address or others.
- Subject's public key - an asymmetric public key that is used by other parties to verify that they are communicating with the subject who is expected to be in control of the corresponding private key. The public key can also be used to encrypt messages that only the subject can decrypt.
- Issuer - an entity that is responsible for validating the identity of the subject. It is usually a [Certificate Authority \(CA\)](#) that is trusted by the consumer of the certificate, but in the case of a self-signed certificate, it can be the subject itself
- Issuer's signature - a signature of the certificate's contents using the issuer's private key

Trusted CAs that serve as the root of trust are usually included in the operating system or a Web browser. This allows applications to verify the certificates of the servers they communicate with. Web browsers use [HyperText Transfer Protocol Secure \(HTTPS\)](#) [Res00] - a variant of [HTTP](#) that uses TLS to secure the underlying [TCP](#) or [QUIC](#) connections of Web applications. The one-way approach where only the server has to authenticate itself to the clients is usually sufficient for most web interactions. TLS can also be used for mutual authentication, where both of the communicating parties have to present a valid certificate to each other, but this requires additional infrastructure to manage the client-side certificates. This mode of operations is sometimes used in Zero Trust networking, in microservice architectures and TLS-based [VPN](#) applications.

TLS is rather complex because it needs to support many possible use cases while remaining backward compatible. It allows the negotiation of security parameters like cipher suits.

TLS operates at the Transport layer (L4) and above so it encrypts the application traffic, but not the IP datagrams. While an [ISP](#) or an intruder with access to the network cannot decrypt the traffic to see what is being communicated, it could see which IP addresses are communicating with each other.

[Secure Shell Protocol \(SSH\)](#) is an Application layer (L7) protocol that allows a client to securely log in and execute commands on a remote server. It does not rely on TLS but uses similar cryptographic primitives. Both the client and the server must authenticate to each other using public-key cryptography or a password. The cryptographic material must be distributed via a side channel either manually or via automation tools.

2.2.2 Network Layer (L3) Security Protocols

Network layer (L3) security protocols provide host-to-host security in Internet Protocol communications. Their implementations are usually configured in the operating systems of the hosts and are invisible to the higher-layer protocols that depend on the Internet Protocol.

IPSec

Internet Protocol Security (IPSec) is a Network layer (L3) protocol suite that provides host-to-host secure communications between [IP](#) hosts. It has two modes of operation:

- Transport mode - encrypts the payload of an IP datagram, but not the header. This mode secures the traffic between two network hosts, but similarly to [TLS](#), the routers between the hosts can still see the IP addresses of the hosts and the ports they are communicating on
- Tunneling mode - encrypts the entire IP datagram and encapsulates it in a new IP datagram with an unencrypted header. This mode is used for secure comm allows a host to decrypt the encapsulated both the traffic connection between a client and a server that can decrypt the encapsulated IP datagram and forward it to its destination.is often used to securely connect a client to a server that can decrypt the encapsulated IP datagrams and forward them to their destination. The client

IPSec always requires mutual authentication, unlike [TLS](#), where it is optional. Authentication can be achieved either via digital certificates, [Pre-Shared Keys \(PSKs\)](#) or others.

OpenVPN Protocol

-
-
-
-

WireGuard

WireGuard [[Don17](#)] is a more recent protocol with a design informed by lessons learned from IPSec and OpenVPN and a key management approach inspired by [SSH](#). It is a lower-level protocol that focuses on configuration simplicity while network topology, peer discovery and key distribution are left as a responsibility of higher-level systems that use it as a building block. Wireguard is implemented as a Layer 3 overlay over UDP tunnels. WireGuard has both user-space implementations that use a TUN driver and also has direct support built into the Linux Kernel since version 5.6 (May 2020). The kernel implementation allows for better performance because it does not need to copy packets between the kernel and user-space memory.

WireGuard's cryptography is based on the **Noise Protocol Framework**[[Per18](#)]. Noise is another recent effort that applies the ideas of TLS in a simplified way. It serves as a blueprint for designing use-case-specific protocols for establishing secure communication channels based on [Eliptic Curve Diffie-Hellman \(ECDH\)](#) handshake patterns. It powers the end-to-end encryption in messaging applications such as WhatsApp and Signal, and [Virtual Private Network \(VPN\)](#) software such as WireGuard and Nebula.

The snippets below show a minimal set of configuration options that need to be provided for two peers to be able to form secure tunnels with each other.

```

1 # peer1.conf
2 [Interface]
3 Address = 101.0.0.1/32

```

VPN This way the VPN client can hide the IP addresses of the final destinations of the traffic from hosts between itself and the VPN server. Depending on the use case, the VPN server could either facilitate the client's secure access to other resources on the virtual network, or provide anonymity to the VPN client by hiding its IP address from the hosts its communicating with on the public Internet.

The Internet Proto-

```
4 ListenPort = 53063
5 PrivateKey = ePTiXXhHjvAHdWUr8Bimk30n0gh3m241RAzsNOJZDW0=
6
7 [Peer]
8 PublicKey = BSn0ejd1Y3bKuD+Xpg0ZZe0f+Ies/oql0NZxw+S0mkc=
9 AllowedIPs = 101.0.0.2/32
10 Endpoint = peer1.example.com:38133

1 # peer2.conf
2 [Interface]
3 Address = 101.0.0.2/32
4 ListenPort = 38133
5 PrivateKey = sN/d6XUPEVPGSziVgCC0n0ivDK+qAoYC3nxxssQ5Rls=
6
7 [Peer]
8 PublicKey = e/TxvPmrgcc1G4cSH2bHv5J0PRHXKjYxTFoU8r+G93E=
9 AllowedIPs = 101.0.0.1/32
```

Each peer has a public/private key pair that is used for authentication and encryption based on the Noise Protocol Framework [Per18]. The **Address** field specifies the virtual IP address that the local network interface will use, while the **AllowedIPs** field specifies what virtual IP addresses are associated with a peer's public key. A peer's **Endpoint** field specifies the URL at which it can be reached. Only one of the peers must be configured with a reachable endpoint for the other one. In the above example once **peer1** initiates communication with **peer2**, **peer2** will learn the current endpoint of **peer1** and will be able to communicate back with it.

2.3 Multiparty Communication - Barriers and Traversal

The version of the Internet Protocol, that was originally deployed globally (IPv4), uses 32-bit numbers as IP addresses, allowing for around 4 billion unique addresses. Due to the popularity of the Internet, there are many more devices than available IPv4 addresses, which has caused challenges. IPv6 is a newer version of the protocol that uses a larger 128-bit address space which is sufficient for assigning 100 addresses for each atom on Earth. However, its adoption has been slow, as according to Google[Goo23] as of 2023 around 41% of their users access their services over IPv6. Additionally, despite that IPv6 allows for all devices to be addressable on the Internet, for security reasons, most of them would use firewalls to block incoming remote traffic that is not associated with outgoing connections.

A widespread solution to the addressing problem is **Network Address Translation (NAT)**. It allows many devices without globally unique IP addresses to initiate connections to publicly addressable devices on the Internet via a limited number of gateways that must have globally unique IP addresses. A NAT gateway replaces the local source IP address of each outgoing IP datagram with its own public IP address before passing it on to the next link on the way to the destination while maintaining a mapping between the source and destination IPs in a translation table. The destination host can then address its responses back to the NAT gateway's public IP address, which in turn replaces its own IP from the incoming datagrams with the IP of the local device and forwards them to it. If the IP datagrams encapsulate TCP/UDP packets, the gateway additionally rewrites the source and destination ports, which means that NAT techniques can be placed somewhere between Layers 3 and 4 of the OSI

CIDR is a common notation for describing IP address ranges, e.g. '192.168.0.1/16' where the number after the slash describes the bit-length of the fixed prefix for a subnet.

model.

The effect of NAT on connectivity is similar to an IPv6 firewall as they both allow devices on a local network to initiate bidirectional communication to remote devices with public IP addresses, but connections cannot be natively initiated by the remote devices. As Figure 2.2 shows, it follows that when two devices are behind separate NATs, neither can contact the other first. **Client/Server** communication is less affected by this limitation because Servers are usually deployed to a public IP address that can be contacted by Clients with local IP addresses. **Peer-to-Peer** communication, however, is more challenging because the peers are often devices in separate residential networks behind different NATs. Several **NAT traversal** techniques try to solve this with different performance tradeoffs and success that varies depending on the NAT [JA07] and its behavior when mapping ports and IP addresses.




Figure 2.2: Two parties behind separate NATs

One approach based on the Client/Server model is to use a publicly addressable **relay** server that is contacted by the NATed devices and then forwards the Peer-to-Peer traffic to the intended recipient. Compared to direct communication, relaying results in a higher network latency due to the longer path that each packet must travel. Maintaining a relay server requires some technical expertise and may be costly depending on the expected throughput. Despite the drawbacks, relaying works under most networking scenarios and is therefore often used as a fallback in case all other approaches fail to find a direct path. Protocols such as **Traversal Using Relays around NAT (TURN)** [Red+20] and **Designated Encrypted Relay for Packets (DERP)** [Tai22] can be used to securely implement relaying.

The NAT gateway in many residential networks is a Router device under the customer's control that has a statically or dynamically assigned public IP address. Most routers can be manually configured through their admin page to forward all traffic that arrives at a given port to a specific device on the local network. Remote applications can then initiate a connection to the local device if they know the IP address of the router and the forwarded port. The manual configuration, however, can be inconvenient and many users may be unaware of that setting because it is not necessary for the more straightforward Client/Server communications. Some routers also support programmatic configuration of port forwarding via a Layer 7 protocol like **Universal Plug and Play (UPnP)** or its successors **NAT Port Mapping Protocol (NAT-PMP)** and **Port Control Protocol (PCP)**. However, these protocols are not always supported and are often disabled by the local network administrators due to security concerns related to bugs

describe some of the nat behaviors, e.g. if the source IP address/-port vary per destination are changed depending on the destination/port mapping algorithms, if it maps ports, IPs, whether the mapped IPs are different per destination and others.

in their implementation, vulnerable IOT devices on the local network or malicious programs being able to expose local devices to the internet.

An efficient NAT traversal approach that works with some types of NATs is to use [Session Traversal Utilities for NAT \(STUN\)](#) [Pet+20] in combination with [UDP](#) hole punching (Figure 2.3). STUN is a protocol operating at Layer 7 that allows a client application to detect the presence of NAT gateways on the network path to a public STUN server, and identify their types and the public IP address that they map to externally. The process usually involves the following steps:

1. An application sends [UDP](#) datagrams to the STUN server
2. The STUN server responds with the source IP address and port specified inside the datagrams
3. The application compares its own endpoint with the source endpoint observed by the STUN server and if the values differ, it can be inferred that they were rewritten by a NAT. Additional STUN servers are contacted to determine if the NAT maps IPs and ports in a predictable fashion.

UDP hole punching is a related technique that, depending on the NAT types, can allow direct communication between two applications behind separate NATs. The applications must discover each other's externally mapped endpoints, perhaps via the STUN server. If the NATs use the same external port regardless of the remote destination, the two applications can:

1. simultaneously send UDP packets to each other's external endpoints.
2. their respective NATs will see the outgoing connection to the other peer and create a port mapping for the reverse connection - a "punched hole"
3. when the incoming traffic arrives from it and forward it correctly.

NATs that map different ports per remote destination sometimes allocate port numbers predictably, which can be used by the peers to try to guess the port that will be opened by the opposing side's NAT.




Figure 2.3: NAT traversal via STUN

In mobile networks like 4G and 5G, the [ISP](#) often utilizes a [Carrier-Grade NAT \(CGNAT\)](#) as part of their infrastructure, while all devices under the user's control, including the router, only have local IP addresses. STUN techniques would fail to discover a direct path between two

parties behind separate CGNATs or other unpredictable NAT algorithms. The only remaining possibility is to relay the traffic via a publicly reachable third-party host using a protocol similar to TURN.

Interactive Connectivity Establishment (ICE) is a protocol that describes a standard way for peers to gather candidate addresses for direct communication via STUN and TURN and then exchange them via a signaling server. The protocol continuously checks which candidates provide the best connection and adjusts them.

Web Real-Time Communication (WebRTC) is a framework that allows peer-to-peer communications between Web applications in Web browsers. Web applications are normally limited to HTTP connections and cannot use raw UDP or TCP connections. WebRTC implements the ICE functionality in Web browsers and provides an API to Web applications.

only 65000 ports per IP address means that CGNATs that provide more than 65000 connections from client devices require more than one public IP address

hairpinning - Hairpinning, also known as NAT loopback or NAT reflection, is a technique used by NAT devices to allow hosts on a private network to access a public server using its public IP address. Without hairpinning, the NAT device would not recognize the connection as a loopback connection and would route it to the public network, causing the con-

Chapter 3

Overlay Networks

An **overlay network** is a higher-order solution that provides additional networking functionality on top of an existing underlay network like the Internet. From the point of view of its consumers, an overlay network may appear at a lower OSI layer, despite being implemented using protocols from higher layers. For example [VPNs](#) can provide virtual interfaces to the Operating System at the Link layer (L2) or Network layer (L3) while being implemented on top of a Transport layer (L4) protocol like [User Datagram Protocol \(UDP\)](#) or a Presentation layer (L6) protocol like [Transport Layer Security \(TLS\)](#). Virtual IP addresses can be assigned to the hosts and applications that are already designed to work with TCP/IP can directly use the virtual network interfaces via the regular TCP/IP mechanisms provided by the operating system.

Other overlay networks are both implemented and used at the Application layer (L7). To communicate via such an overlay network, applications often have to implement specific functionality in their software by utilizing a framework or a library.

Figure [3.1](#) shows an approximate OSI model mapping of several protocols and network overlay solutions from the point of view of the systems that use them and the arrows show dependency relations between them.




Figure 3.1: OSI model mapping of various protocols

3.1 Network layer (L3) and Data Link (L2) Overlay Networks

3.1.1 Traditional VPNs

Layer 2 virtual networks provide a virtual network switch that allows remote machines to be on the same virtual LAN and share the same IP address range. Layer 3 virtual networks provide a virtual network router that allows remote machines to be on separate LANs. Depending on the specific implementation, the overlay network can either be implemented directly in the Operating System's kernel, or on top of a driver like **TUN** or **TAP**.

- Layer 2 vs Layer 3 Networks
 - Layer 2 overlay networks bridge other networks
 - * virtual network switch
 - * remote machines are on the same virtual LAN and can share the same IP address range
 - * allows broadcast/multicast
 - * TAP driver
 - Layer 3 overlays route traffic between separate local networks
 - * virtual network router
 - * remote machines are on separate LANs
 - * simpler to configure
 - * TUN driver
 - * the low-level solutions from the previous section are complex to set up.
 - * overlay networks package some of those solutions for a specific use case. Most overlay networks use a combination of the NAT traversal techniques mentioned previously. They can be placed in Layers 2, 3 or 7. Layer 2 overlays act as a virtual network switch, while Layer 3 overlays act as a virtual network router. Layer 7 overlays are implemented in user-space as libraries or applications that run on top of the network stack of the host operating system. Layer 2 and 3 overlays can either be implemented as kernel modules or as user-space applications that use a **TUN/TAP** driver to interface with the kernel.
 - * The term “VPN” is somewhat overloaded as it can refer to different related concepts.

VPNs are implemented as Layer 2 or 3 network overlays. They are commonly used for securely connecting machines from different **LANs**. They provide software emulation of a network interface controller via a TUN/TAP driver on the operating system level and allow other software to transparently use the functionality of the **Internet Protocol (IP)** suite without requiring extra changes. Traditional **VPNs** such as IPsec [FK11] and OpenVPN [Ope22] use a centralized service that all (encrypted) client communications must pass through. This introduces a single point of failure and a potential bottleneck that might negatively impact the performance of the multiparty computations due to their **Peer to Peer (P2P)** nature.

explain the names

tls vs ipsec vpns. TLS vpns offer a virtual network interface at layer 3, but run over L7 TLS

OpenVPN Client/Server

We already discussed the VPN protocol behind OpenVPN. Here we will go into the client and server software and how they are implemented and discuss the typical network topologies.

3.1.2 Mesh VPNs

- Tinc
- N2N
- Tailscale
- Nebula
- ZeroTier

Mesh VPNs such as Tinc [Sli22], Tailscale [Tai] and Nebula [Def22] utilize NAT Traversal techniques to create direct P2P links between the clients for the data traffic. Authentication, authorization and traffic encryption are performed using certificates based on public key cryptography.

All three are open-source, except Tailscale's coordination service which handles peer discovery and identity management. Headscale [Fon22] is a community-driven open-source alternative for that component. Tinc is the oldest of the three but has a relatively small community. It is mainly developed by a single author and appears to be more academic than industry motivated. Nebula and Tailscale are both business driven. Tailscale was started by some high-profile ex-googlers and is the most end-user-focused of the three, providing a service that allows people to sign up using identity providers such as Google, Microsoft, GitHub and others. They also provide an Admin console that allows a user to easily add their personal devices to a network or share them with others. It also has support for automation tools like Terraform for creating authorization keys and managing an Access Control List (ACL) based firewall. Nebula was originally developed at the instant messaging company Slack to create overlay networks for their cross-region cloud infrastructure, but the authors later started a new company and are currently developing a user-centric platform similar to Tailscale's. Nebula is more customizable than Tailscale and since it is completely open-source it can be adapted to different use cases, but it is also more involved to set up. A certificate authority needs to be configured for issuing the identities of the participating hosts. Furthermore, publicly accessible coordination servers need to be deployed to facilitate the host discovery. Tailscale employs a distributed relay network of Designated Encrypted Relay for Packets (DERP) servers, while Nebula can be configured to route via one of the other peers in the VPN.

3.2 Application Layer (L7) Network Overlays

3.2.1 OpenZiti

- uses relays
- libP2P
- ngrok
- TOR
- BitTorrent
- IPFS
- Ethereum
- Teleport

- Freenet

Part II

Proof of Concepts and Evaluation

Chapter 4

Testing methodology

In the following chapters, we will design and implement several solutions for ad hoc MPC sessions based on a subset of the previously discussed related works:

- Internet protocol
- Wireguard
- Tailscale
- Headscale
- ? Headscale with DID identity?
- ? WebRTC?
- A custom solution that automates the WireGuard configuration by visiting a web page

Additionally, we will analyze and compare them in terms of performance, security and usability

4.1 Measuring performance

During the preparation phase of the project, we developed the [Extensible Evaluation Environment \(\$E^3\$ \)](#) framework which simplifies and automates the process of deploying machines in different geographical regions, connecting them via an overlay network and executing multiparty computations between them, where each machine represents a different party.

To summarize, E^3 is a set of scripts that use several automation tools:

- Terraform - declarative provisioning
- NixOS - declarative Linux distribution
- Colmena - declarative deployment for NixOS
- PSSH - parallel execution of remote scripts over ssh
- DigitalOcean - a cloud provider

It allows us to quickly provision cloud virtual machines in multiple regions and reproducibly deploy all necessary software for running a multiparty computation over a chosen network overlay solution. The source code of E^3 can be found on [GitHub](#)

Each solution will be deployed using the E^3 framework and the performance will be quantitatively measured in terms of the time it takes to execute several MPyC demos. The selected

demos have different complexities in terms of communication rounds and message sizes which will allow us to observe their impact on the overall performance.

1. Secret Santa - high round complexity with small messages
2. Convolutional Neural Network (CNN) MNIST classifier - low round complexity with large messages

The demos will be configured at three different input size levels

- Low,
- Medium
- High

Furthermore, the demos will be executed in several networking scenarios:

1. 1-10 parties in the same geographic region
2. 1-10 parties evenly distributed across two nearby regions
3. 1-10 parties evenly distributed across two distant regions
4. 1-10 parties distributed across multiple distant regions

4.2 Security

We will analyze aspects such as

- key distribution
- trust model - are there any trusted third parties and what would be the consequences if they are corrupted or breached
- traffic encryption
- identity strength

4.3 Usability

For each solution, we will describe the steps that the parties need to perform to execute a joint multiparty computation. Those steps will be analyzed in terms of:

- Complexity - how much technical expertise is expected from the parties to be able to execute the steps
- Initial effort - how much effort is each party expected to put in preparing for their first joint computation
- Repeated effort - after the initial setup, how much effort is required to perform another computation
 - with the same set of parties
 - with another set of parties
- Finalization effort - how much effort is required to finalize the MPC session once it is complete and clean up any left-over artifacts or resources so that the machine of each party is in its original state

Chapter 5

Internet Protocol based solution

This solution focuses on directly using the internet protocol without involving an overlay network. Our goal is to analyze the implications of using only the functionalities that MPyC directly supports to serve as the reference for our later experiments.

5.1 Implementation details

We will manually set up the multiparty computations via the public IP addresses of the machines and DNS.

5.2 Performance analysis

5.3 Security analysis

5.4 Usability analysis

Chapter 6

WireGuard based solution

This solution creates an overlay network by manually configuring WireGuard on each machine.

6.1 Implementation details

6.2 Performance analysis

6.3 Security analysis

6.4 Usability analysis

Chapter 7

Tailscale based solution

Tailscale is a VPN solution that configures a mesh of direct Wireguard tunnels between the peers.

7.1 Implementation details

7.2 Performance analysis

7.3 Security analysis

7.3.1 Trust model

There is a centralized service that deals with the key distribution, which needs to be trusted to provide the correct public keys for the correct parties

7.3.2 Identity

Identity is based on third party identity providers such as Microsoft and GitHub

- Magic DNS
-

7.4 Usability analysis

With tailscale each party needs to

- register a Tailscale account
- Download and install tailscale on the machine they want to run a multiparty computation
- Run tailscale on their machine and logs into their account in order to link it to their own Tailnet
- Share their Tailscale machine with the Tailnets of each of the other parties
- Download the demo they want to run
- Form the flags for running the chosen demo

- add -P \$HOST:\$PORT for each party using their Tailscale hostname/virtual IP
- Run the demo

Chapter 8

Headscale based solution

This solution is similar to the Tailscale one, but it uses Headscale - a self-hosted open-source alternative to the closed-source Tailscale coordination service.

8.1 Implementation details

8.2 Performance analysis

8.3 Security analysis

8.3.1 Trust model

There still is a centralized service like in the Tailscale solution, but here it is self-deployed.

8.3.2 Identity

8.4 Usability analysis

Bibliography

- [81] *Internet Protocol*. Request for Comments RFC 791. Internet Engineering Task Force, Sept. 1981. 51 pp. DOI: [10.17487/RFC0791](https://doi.org/10.17487/RFC0791). URL: <https://datatracker.ietf.org/doc/rfc791> (visited on 05/14/2023) (cit. on p. 8).
- [Def22] Defined. *Nebula: Open Source Overlay Networking | Nebula Docs*. 2022. URL: <https://docs.defined.net/docs/> (visited on 12/01/2022) (cit. on p. 20).
- [Don17] Jason A. Donenfeld. “WireGuard: Next Generation Kernel Network Tunnel”. In: *Proceedings 2017 Network and Distributed System Security Symposium*. Network and Distributed System Security Symposium. San Diego, CA: Internet Society, 2017. ISBN: 978-1-891562-46-4. DOI: [10.14722/ndss.2017.23160](https://doi.org/10.14722/ndss.2017.23160). URL: <https://www.ndss-symposium.org/ndss2017/ndss-2017-programme/wireguard-next-generation-kernel-network-tunnel/> (visited on 09/28/2022) (cit. on p. 11).
- [FK11] Sheila Frankel and Suresh Krishnan. *IP Security (IPsec) and Internet Key Exchange (IKE) Document Roadmap*. Request for Comments RFC 6071. Internet Engineering Task Force, Feb. 2011. 63 pp. DOI: [10.17487/RFC6071](https://doi.org/10.17487/RFC6071). URL: <https://datatracker.ietf.org/doc/rfc6071> (visited on 05/28/2023) (cit. on p. 19).
- [Fon22] Juan Font. *Juanfont/Headscale*. Dec. 6, 2022. URL: <https://github.com/juanfont/headscale> (visited on 12/06/2022) (cit. on p. 20).
- [Goo23] Google. *IPv6 – Google*. May 21, 2023. URL: <https://www.google.com/intl/en/ipv6/statistics.html> (visited on 05/21/2023) (cit. on p. 12).
- [JA07] Cullen Fluffy Jennings and Francois Audet. *Network Address Translation (NAT) Behavioral Requirements for Unicast UDP*. Request for Comments RFC 4787. Internet Engineering Task Force, Jan. 2007. 29 pp. DOI: [10.17487/RFC4787](https://doi.org/10.17487/RFC4787). URL: <https://datatracker.ietf.org/doc/rfc4787> (visited on 05/26/2023) (cit. on p. 13).
- [LY06] Chris M. Lonvick and Tatu Ylonen. *The Secure Shell (SSH) Protocol Architecture*. Request for Comments RFC 4251. Internet Engineering Task Force, Jan. 2006. 30 pp. DOI: [10.17487/RFC4251](https://doi.org/10.17487/RFC4251). URL: <https://datatracker.ietf.org/doc/rfc4251> (visited on 05/28/2023) (cit. on p. 8).
- [Ope22] OpenVPN. *Community Resources*. OpenVPN. 2022. URL: <https://openvpn.net/community-resources/> (visited on 11/30/2022) (cit. on p. 19).
- [Per18] Trevor Perrin. “The Noise Protocol Framework”. In: (July 1, 2018) (cit. on pp. 11, 12).

- [Pet+20] Marc Petit-Huguenin et al. *Session Traversal Utilities for NAT (STUN)*. Request for Comments RFC 8489. Internet Engineering Task Force, Feb. 2020. 67 pp. DOI: [10.17487/RFC8489](https://datatracker.ietf.org/doc/rfc8489). URL: <https://datatracker.ietf.org/doc/rfc8489> (visited on 05/26/2023) (cit. on p. 14).
- [Red+20] Tirumaleswar Reddy.K et al. *Traversal Using Relays around NAT (TURN): Relay Extensions to Session Traversal Utilities for NAT (STUN)*. Request for Comments RFC 8656. Internet Engineering Task Force, Feb. 2020. 79 pp. DOI: [10.17487/RFC8656](https://datatracker.ietf.org/doc/rfc8656). URL: <https://datatracker.ietf.org/doc/rfc8656> (visited on 04/24/2023) (cit. on p. 13).
- [Res00] Eric Rescorla. *HTTP Over TLS*. Request for Comments RFC 2818. Internet Engineering Task Force, May 2000. 7 pp. DOI: [10.17487/RFC2818](https://datatracker.ietf.org/doc/rfc2818). URL: <https://datatracker.ietf.org/doc/rfc2818> (visited on 05/26/2023) (cit. on p. 10).
- [Res18] Eric Rescorla. *The Transport Layer Security (TLS) Protocol Version 1.3*. Request for Comments RFC 8446. Internet Engineering Task Force, Aug. 2018. 160 pp. DOI: [10.17487/RFC8446](https://datatracker.ietf.org/doc/rfc8446). URL: <https://datatracker.ietf.org/doc/rfc8446> (visited on 05/26/2023) (cit. on p. 9).
- [Sli22] Guus Sliepen. *Tinc Docs*. Nov. 30, 2022. URL: <https://www.tinc-vpn.org/docs/> (visited on 11/30/2022) (cit. on p. 20).
- [Tai] Tailscale. *Tailscale*. Tailscale. URL: <https://tailscale.com/kb/> (visited on 11/30/2022) (cit. on p. 20).
- [Tai22] Tailscale. *DERP Servers*. Tailscale. Dec. 21, 2022. URL: <https://tailscale.com/kb/1232/derp-servers/> (visited on 05/28/2023) (cit. on p. 13).