Secure Sessions for Ad Hoc Multiparty Computation in MPyC

Master thesis

Emil Nikolov

Id nr: 0972305
`emil.e.nikolov@gmail.com`

Supervisor : Dr. ir. L.A.M. (Berry) Schoenmakers

April 25, 2023

Contents

List of Abbreviations

UDP User Datagram Protocol. 3, 6

UPnP Universal Plug and Play. 5

VPN Virtual Private Network. 4, 9, 10

List of Figures

Chapter 1

Introduction

Chapter 2

Related work

In this chapter, we will examine prior works relevant to our problem of connecting multiple parties over The Internet for a joint MPyC computation. To identify the challenges, we will begin with an overview of The Internet Protocol Suite, which is central to the modern Internet. Following that, we will review several existing network overlay solutions that we can later use as building blocks for an overall solution to our problem. To maintain a consistent mental picture of the solutions, we will map them to the layers of the Open Systems Interconnection (OSI) model. While many protocols implement aspects of several layers and do not strictly fit inside the OSI model, it is still a useful visualization tool.

2.1 The Internet

The modern Internet is a global multi-tiered network of devices that communicate using the protocols of the Internet Protocol Suite (TCP/IP). Typically, an Internet Service Provider (ISP) manages the physical infrastructure that connects an end-user's devices with the rest of the internet. Figure 2.1 describes the responsibilities of the 7 layers of the OSI model and how they fit with the protocols of the Internet Protocol Suite, which are often referred to as the TCP/IP model. The newer TCP/IP model only recognizes 4 layers as it merges the OSI Session (L5) and Presentation (L6) layers into the Application layer (L7), as well as the Physical layer (L1) into the Data link layer (L2). While the TCP/IP model is a more accurate representation of The Internet, the 7-layer numbers of the OSI model are still widely used in the literature.

The **Internet Protocol (IP)** is a Network (L3) protocol of the Internet Protocol Suite that is responsible for delivering datagrams between devices across the boundaries of their Local Area Networks (LANs) by possibly routing traffic via multiple intermediate devices (routers). A datagram is a connectionless packet that is delivered on a best-effort basis. IP datagrams have a header that contains fields such as the **IP addresses** of its source and destination, and a payload that encapsulates the data from the protocols of the layers above. Routers are devices that are part of multiple networks and relay datagrams between them based on a routing table that maps IP address ranges (Classless Inter-Domain Routing (CIDR)) to networks.

show a diagram of the internet with multiple local networks

User Datagram Protocol (UDP) and **Transmission Control Protocol (TCP)** are Transport (L4) protocols that add the concept of ports to allow having multiple communication
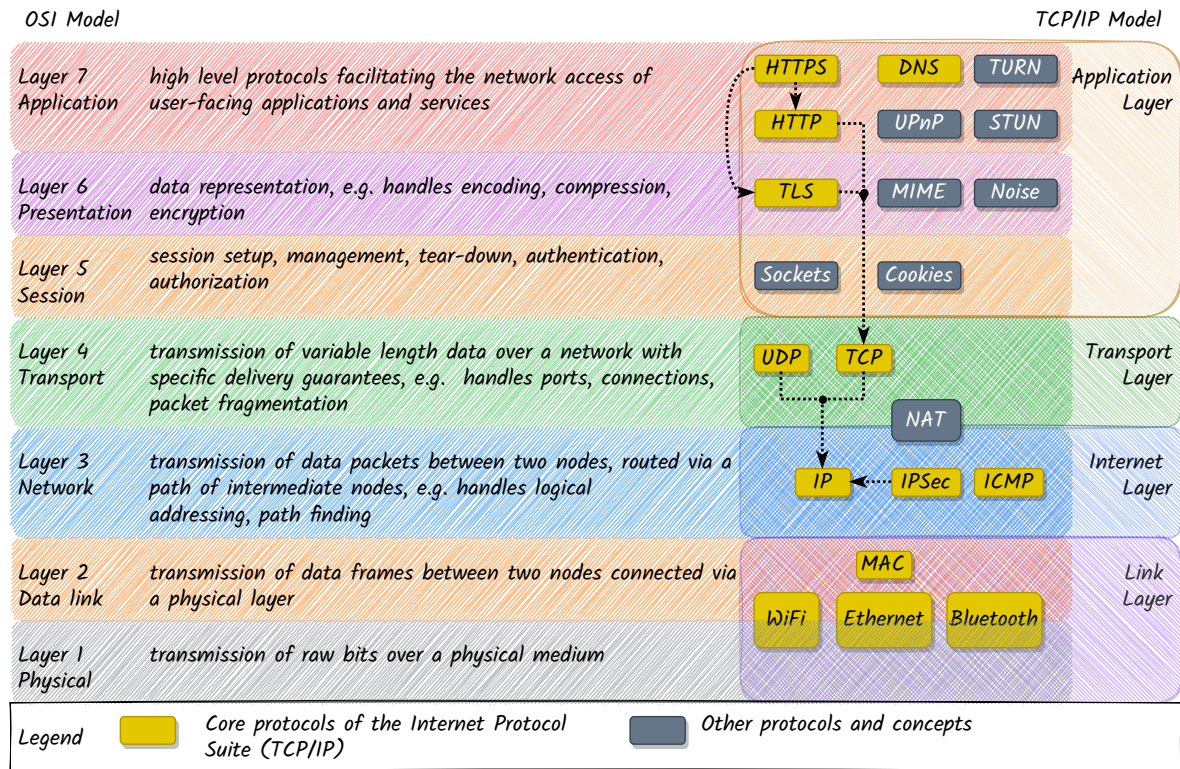
Figure 2.1: OSI model mapping of the Internet Protocol Suite

channels simultaneously between two devices. UDP provides best-effort delivery, while TCP is a reliable transport with delivery guarantees. TCP maintains stateful connections and handles acknowledgments and retransmissions in case packets are lost in transit.

Transport Layer Security (TLS) is a protocol that adds encryption on top of a reliable transport protocol such as TCP. It is usually placed in the Presentation Layer (L6), but it does not strictly fit in any single OSI layer. It is rather complex because it needs to support many possible use cases across the internet. The **Noise Protocol Framework** [Per18] is a more recent effort that applies the ideas of TLS in a simplified way by serving as a blueprint for designing use-case specific protocols for establishing secure communication channels based on Eliptic Curve Diffie-Hellman (ECDH) handshake patterns. It powers the end-to-end encryption in messaging applications such as WhatsApp and Signal, and Virtual Private Network (VPN) software such as WireGuard and Nebula.

tls noise is transport agnostic noise has limited cipher suites

The version of the Internet Protocol that was originally deployed globally (IPv4) uses 32-bit numbers as IP addresses, allowing for around 4 billion unique addresses. Due to the popularity of the internet, there are many more devices than available IPv4 addresses, which has caused challenges. IPv6 is a newer version of the protocol that uses a larger 128-bit address space which is sufficient for assigning 100 addresses for each atom on Earth. Its adoption has been slow, as according to Google as of 2023 only around 37% of the requests to their services use IPv6. Additionally, even though IPv6 allows for each device to be addressable on the Internet, most of them would still use firewalls to block incoming traffic that is not associated with outgoing connections.

A more widespread solution to the addressing problem is Network Address Translation (NAT). It allows many devices without globally unique IP addresses to initiate connections to publicly addressable devices on the Internet via a limited number of gateways that must have globally unique IP addresses. A NAT gateway replaces the local source IP address of each outgoing IP datagram with its own public IP address before passing it on to the next link on the way to the destination while maintaining a mapping between the source and destination IPs in a translation table. The destination host can then address its responses back to the NAT gateway's public IP address, which in turn replaces its own IP from the incoming datagrams with the IP of the local device and relays them to it. If the IP datagrams encapsulate TCP/UDP packets, the gateway additionally rewrites the source and destination ports, which means that NAT techniques can be placed somewhere between Layers 3 and 4 of the OSI model.

The NAT approach allows devices with local IP addresses to initiate bidirectional connections with remote devices across the public internet but it does not natively allow those remote devices to initiate a connection first. Furthermore, as Figure 2.2 shows, when two different devices are behind separate NATs, neither can contact the other first. Client/Server communications are less affected by this limitation because the Server can be deployed to a machine with a public IP address that can be contacted by the Client. Peer-to-Peer communications, however, are more challenging, because multiple peers can be behind separate NATs. Several **NAT traversal** techniques try to solve this with varying success and different performance tradeoffs.

Figure 2.2: Two parties behind separate NATs

As in the Client/Server model, it is possible to relay the peer-to-peer communications via another machine with a public IP address. While the additional hop on the network path makes this approach less efficient than direct communication, it works under most networking scenarios and is therefore often used as a fallback in case other approaches fail to find a direct path. Protocols such as Traversal Using Relays around NAT (TURN) [**turnRFC**] and Designated Encrypted Relay for Packets (DERP) [**derpDocs**] can be used to securely implement relaying.

In many home networks, the role of a NAT gateway is played by a publicly addressable router under the customer's control, which often allows for port forwarding to be configured either manually or programmatically via a Layer 7 protocol like Universal Plug and Play (UPnP) or its successors NAT Port Mapping Protocol (NAT-PMP) and Port Control Protocol

(PCP). This would allow remote hosts to initiate connections with the local ones if they know the router's public IP address. Those protocols however are not always supported by the router and are often disabled by the local network administrators due to the security concerns of malicious programs being able to request and modify the port mappings in the router's translation tables.

Another approach for NAT traversal that usually achieves good performance, but only works with some types of NATs is to use **Session Traversal Utilities for NAT (STUN)** in combination with UDP hole punching. STUN is a set of Client/Server protocols that allow a device to discover information about the NAT gateways that might connect it to The Internet. This includes the IP address of a NAT gateway and its behavior such as if it maps ports, IPs, whether the mapped IPs are different per destination and others. To achieve this, the device uses UDP to contact a public STUN server which responds with the source IP address and port of the incoming datagrams. If a NAT gateway rewrote the IP address and port, then the device will know this by comparing them to its local IP address and port. Additional STUN servers may be contacted to check if the NAT gateway maps consistent IP addresses and ports. UDP hole punching is a related technique that allows two devices

to contact detect its own NAT type and the public address of the NAT gateway via, which relies on the predictable port mapping algorithms that many routers use and a public third party host that can be contacted by the local devices and later serve as an introduction point for them (Figure 2.3).

User Datagram Protocol (UDP) hole punching based on concepts from STUN. The machines of each party can contact a public STUN server 2.3, which will note what Internet Protocol (IP) addresses the connections come from and inform the parties. Since the parties initiated the connection to the STUN server, their routers will keep a mapping between their local IP addresses and the port that was allocated for the connection to be able to forward the incoming traffic. Those "holes" in the NATs were originally intended for the STUN server, but mesh VPNs use the stateless "fire and forget" UDP protocol for their internal communication, which does not require nor provides a mechanism for the NATs to verify who sent a UDP packet. With most NATs, this is enough to be able to (ab)use the "punched holes" for the purpose of Peer to Peer (P2P) traffic from other parties. Mesh VPNs implement the stateful Transmission Control Protocol (TCP) and Transport Layer Security (TLS) protocols on top of UDP and expose a regular network interface to the other programs, keeping them shielded from the underlying complexities. Other NAT implementations such as Symmetric NAT and Carrier-Grade NATs (CGNATs) can be more difficult to "punch through" due to their more complex port mapping strategies. In those cases, establishing P2P connections might involve guesswork or even fail and require falling back to routing the (encrypted) traffic via another party or service.

In most mobile networks (4G, 5G) the ISP utilizes a **CGNAT** as part of their infrastructure, while all devices under the user's control, including the router, only have local IP addresses. STUN techniques would fail to discover a direct path between two parties behind separate CGNATs or other unpredictable NAT algorithms. The only remaining possibility is for a like

Figure 2.3: NAT traversal via STUN

Traversal Using Relays around NAT (TURN), where a publicly reachable third-party host is used not only for introducing the peers but also for relaying all (possibly encrypted) traffic between them.

- Hairpinning

- **Internet Protocol Security (IPSec)**

    - Layer 3 protocol suite part of the Internet Protocol Suite
    - used inside VPN software
    - has implementations in both user and kernel space as well as hardware implementations
    - rewrites and encrypts the IP headers and payloads
    - virtual routing table
    - Initially was built into IPv6, separate from IPv4

2.2   Network overlays

Most Network Overlay solutions can be placed in Layers 2, 3 or 7. Layer 2 overlays act as a virtual network switch, while Layer 3 overlays act as a virtual network router. Layer 7 overlays are implemented in user-space as libraries or applications that run on top of the network stack of the host operating system. Layer 2 and 3 overlays can either be implemented as kernel modules or as user-space applications that use a **TUN/TAP** driver to interface with the kernel.

Figure 2.4 shows an approximate OSI model mapping of several protocols and network overlay solutions from the point of view of the systems that use them and the arrows show dependency relations between them.

2.2.1   TUN/TAP driver

- Layer 2 vs Layer 3 Networks
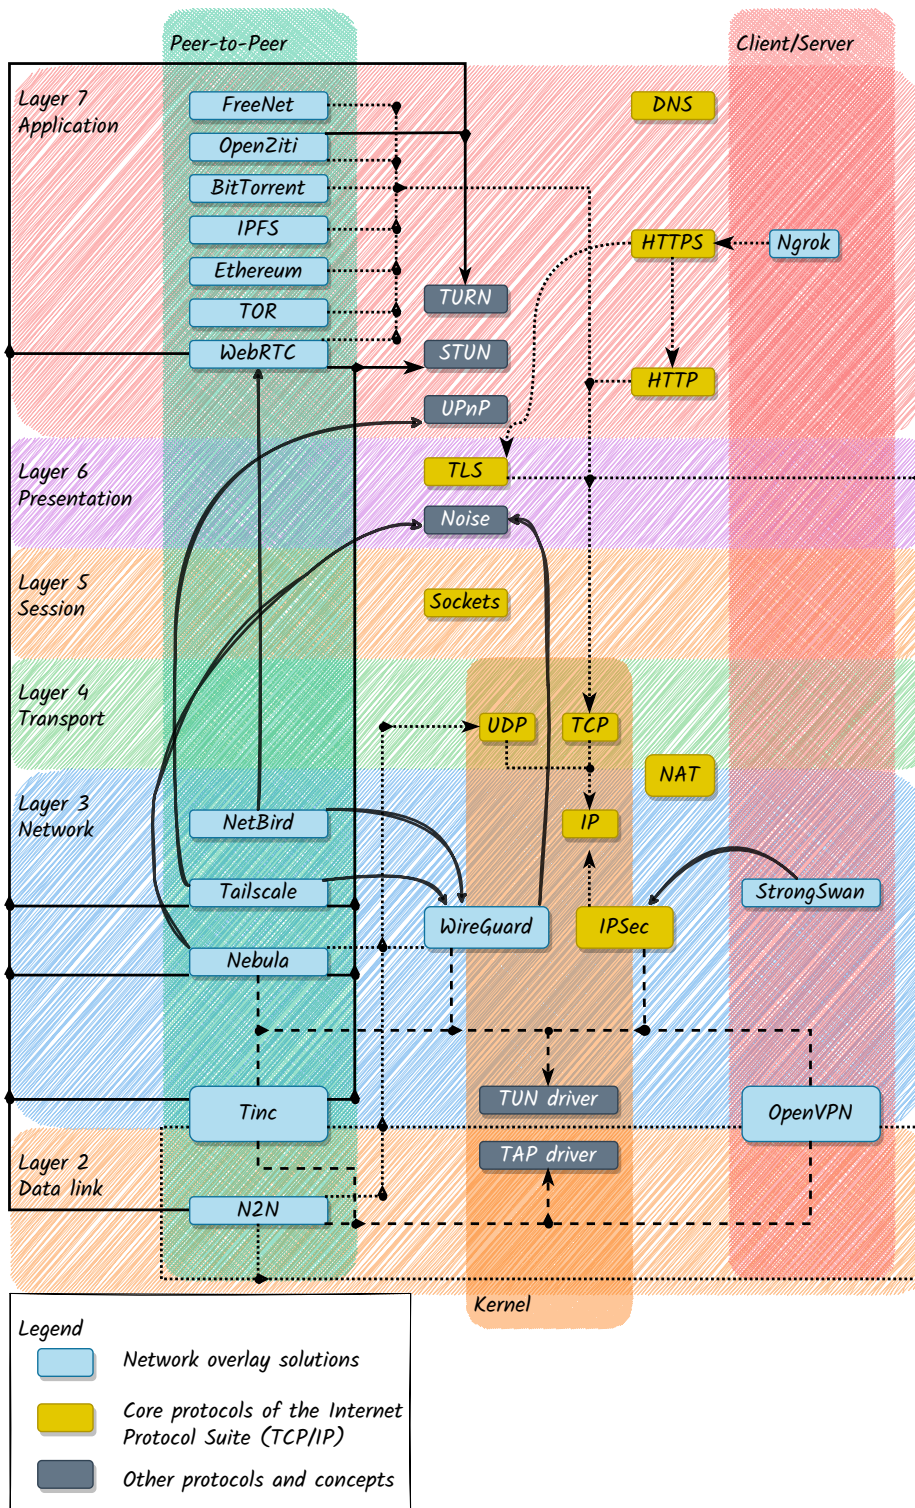   - Layer 2 overlays bridge networks
        * virtual network switch

Figure 2.4: OSI model mapping of various protocols

* remote machines are on the same virtual LAN and can share the same IP address range
* allows broadcast/multicast
* TAP driver
  - Layer 3 overlays route traffic between separate local networks
    * virtual network router
    * remote machines are on separate LANs
    * simpler to configure
    * TUN driver

2.2.2 Traditional VPNs

VPNs are implemented as Layer 2 or 3 network overlays. They are commonly used for securely connecting machines from different LANs. They provide software emulation of a network interface controller via a TUN/TAP driver on the operating system level and allow other software to transparently use the functionality of the IP suite without requiring extra changes. Traditional VPNs such as IPSec [**ipSecDocs**] and OpenVPN [Ope22] use a centralized service that all (encrypted) client communications must pass through. This introduces a single point of failure and a potential bottleneck that might negatively impact the performance of the multi-party computations due to their P2P nature.

2.2.3 WireGuard

WireGuard [Don17] is a more recent protocol with a design informed by lessons learned from IPSec and OpenVPN and a key management approach inspired by SSH. It is a lower-level protocol that focuses on configuration simplicity while network topology, peer discovery and key distribution are left as a responsibility of higher-level systems that use it as a building block. Wireguard is implemented as a Layer 3 overlay over UDP tunnels. WireGuard has both user-space implementations that use a TUN driver and also has direct support built into the Linux Kernel since version 5.6 (May 2020). The kernel implementation allows for better performance because it does not need to copy packets between the kernel and user-space memory.

The snippets below show a minimal set of configuration options that need to be provided for two peers to be able to form secure tunnels with each other.

```
1   # peer1.conf
2   [Interface]
3   Address = 101.0.0.1/32
4   ListenPort = 53063
5   PrivateKey = ePTiXXhHjvAHdWUr8Bimk3On0gh3m241RAzsN0JZDW0=
6
7   [Peer]
8   PublicKey = BSn0ejd1Y3bKuD+Xpg0ZZeOf+Ies/oql0NZxw+SOmkc=
9   AllowedIPs = 101.0.0.2/32
10  Endpoint = peer1.example.com:38133
```

```
1   # peer2.conf
2   [Interface]
3   Address = 101.0.0.2/32
4   ListenPort = 38133
```

```
5   PrivateKey = sN/d6XUPEVPGSziVgCCOnOivDK+qAoYC3nxnssQ5Rls=
6
7   [Peer]
8   PublicKey = e/TxvPmrgcc1G4cSH2bHv5JOPRHXKjYxTFoU8r+G93E=
9   AllowedIPs = 101.0.0.1/32
```

Each peer has a public/private key pair that is used for authentication and encryption based on the Noise Protocol Framework [Per18]. The `Address` field specifies the virtual IP address that the local network interface will use, while the `AllowedIPs` field specifies what virtual IP addresses are associated with a peer's public key. A peer's `Endpoint` field specifies the URL at which it can be reached. Only one of the peers must be configured with a reachable endpoint for the other one. In the above example once `peer1` initiates communication with `peer2`, `peer2` will learn the current endpoint of `peer1` and will be able to communicate back with it.

2.2.4   Mesh VPNs

- Tinc
- N2N
- Tailscale
- Nebula
- ZeroTier

Mesh VPNs such as Tinc [Sli22], Tailscale [Tai] and Nebula [Def22] utilize NAT Traversal techniques to create direct P2P links between the clients for the data traffic. Authentication, authorization and traffic encryption are performed using certificates based on public key cryptography.

All three are open-source, except Tailscale's coordination service which handles peer discovery and identity management. Headscale [Fon22] is a community-driven open-source alternative for that component. Tinc is the oldest of the three but has a relatively small community. It is mainly developed by a single author and appears to be more academic than industry motivated. Nebula and Tailscale are both business driven. Tailscale was started by some high-profile ex-googlers and is the most end-user-focused of the three, providing a service that allows people to sign up using a variety of identity providers including Google, Microsoft, GitHub and others. They also provide an Admin console that allows a user to easily add their personal devices to a network or share them with others. It also has support for automation tools like Terraform for creating authorization keys and managing an Access Control List (ACL) based firewall. Nebula was originally developed at the instant messaging company Slack to create overlay networks for their cross-region cloud infrastructure, but the authors later started a new company and are currently developing a user-centric platform similar to Tailscale's. Nebula is more customizable than Tailscale and since it is completely open-source it can be adapted to different use cases, but it is also more involved to set up. A certificate authority needs to be configured for issuing the identities of the participating hosts. Furthermore, publicly accessible coordination servers need to be deployed to facilitate the host discovery. Tailscale employs a distributed relay network of Designated Encrypted Relay for Packets (DERP) servers, while Nebula can be configured to route via one of the other peers in the VPN.

2.2.5 Layer 7 overlays

WebRTC is

- WebRTC
  - Uses STUN/TURN
  - d
- OpenZiti
  - uses relays
- ngrok
- TOR
- BitTorrent
- IPFS
- Ethereum
- Teleport
- Freenet

Chapter 3

Testing methodology

In the following chapters, we will design and implement several solutions for ad hoc MPC sessions based on a subset of the previously discussed related works:

- Internet protocol
- Wireguard
- Tailscale
- Headscale
- ? Headscale with DID identity?
- ? WebRTC?
- A custom solution that automates the WireGuard configuration by visiting a web page

Additionally, we will analyze and compare them in terms of performance, security and usability

3.1 Measuring performance

During the preparation phase of the project, we developed the Extensible Evaluation Environment (E^3) framework which simplifies and automates the process of deploying machines in different geographical regions, connecting them via an overlay network and executing multiparty computations between them, where each machine represents a different party.

To summarize, E^3 is a set of scripts that use several automation tools:

- Terraform - declarative provisioning
- NixOS - declarative Linux distribution
- Colmena - declarative deployment for NixOS
- PSSH - parallel execution of remote scripts over ssh
- DigitalOcean - a cloud provider

It allows us to quickly provision cloud virtual machines in multiple regions and reproducibly deploy all necessary software for running a multiparty computation over a chosen network overlay solution. The source code of E^3 can be found on GitHub

Each solution will be deployed using the E^3 framework and the performance will be quantitatively measured in terms of the time it takes to execute several MPyC demos. The selected

demos have different complexities in terms of communication rounds and message sizes which will allow us to observe their impact on the overall performance.

1. Secret Santa - high round complexity with small messages
2. Convolutional Neural Network (CNN) MNIST classifier - low round complexity with large messages

The demos will be configured at three different input size levels

- Low,
- Medium
- High

Furthermore, the demos will be executed in several networking scenarios:

1. 1-10 parties in the same geographic region
2. 1-10 parties evenly distributed across two nearby regions
3. 1-10 parties evenly distributed across two distant regions
4. 1-10 parties distributed across multiple distant regions

3.2   Security

We will analyze aspects such as

- key distribution
- trust model - are there any trusted third parties and what would be the consequences if they are corrupted or breached
- traffic encryption
- identity strength

3.3   Usability

For each solution, we will describe the steps that the parties need to perform to execute a joint multiparty computation. Those steps will be analyzed in terms of:

- Complexity - how much technical expertise is expected from the parties to be able to execute the steps
- Initial effort - how much effort is each party expected to put in preparing for their first joint computation
- Repeated effort - after the initial setup, how much effort is required to perform another computation
    - with the same set of parties
    - with another set of parties
- Finalization effort - how much effort is required to finalize the MPC session once it is complete and clean up any left-over artifacts or resources so that the machine of each party is in its original state

Chapter 4

Internet Protocol based solution

This solution focuses on directly using the internet protocol without involving an overlay network. Our goal is to analyze the implications of using only the functionalities that MPyC directly supports to serve as the reference for our later experiments.

4.1 Implementation details

We will manually set up the multiparty computations via the public IP addresses of the machines and DNS.

4.2 Performance analysis

4.3 Security analysis

4.4 Usability analysis

Chapter 5

WireGuard based solution

This solution creates an overlay network by manually configuring WireGuard on each machine.

5.1 Implementation details

5.2 Performance analysis

5.3 Security analysis

5.4 Usability analysis

Chapter 6

Tailscale based solution

Tailscale is a VPN solution that configures a mesh of direct Wireguard tunnels between the peers.

6.1 Implementation details

6.2 Performance analysis

6.3 Security analysis

6.3.1 Trust model

There is a centralized service that deals with the key distribution, which needs to be trusted to provide the correct public keys for the correct parties

6.3.2 Identity

Identity is based on third party identity providers such as Microsoft and GitHub

- Magic DNS

-

6.4 Usability analysis

With tailscale each party needs to

- register a Tailscale account
- Download and install tailscale on the machine they want to run a multiparty computation
- Run tailscale on their machine and logs into their account in order to link it to their own Tailnet
- Share their Tailscale machine with the Tailnets of each of the other parties
- Download the demo they want to run
- Form the flags for running the chosen demo

        – add -P \$HOST:\$PORT for each party using their Tailscale hostname/virtual IP
- Run the demo

Chapter 7

Headscale based solution

This solution is similar to the Tailscale one, but it uses Headscale - a self-hosted open-source alternative to the closed-source Tailscale coordination service.

7.1 Implementation details

7.2 Performance analysis

7.3 Security analysis

7.3.1 Trust model

There still is a centralized service like in the Tailscale solution, but here it is self-deployed.

7.3.2 Identity

7.4 Usability analysis

Bibliography

[Def22]   Defined. *Nebula: Open Source Overlay Networking | Nebula Docs*. 2022. URL: https://docs.defined.net/docs/ (visited on 12/01/2022) (cit. on p. 10).

[Don17]   Jason A. Donenfeld. "WireGuard: Next Generation Kernel Network Tunnel". In: *Proceedings 2017 Network and Distributed System Security Symposium*. Network and Distributed System Security Symposium. San Diego, CA: Internet Society, 2017. ISBN: 978-1-891562-46-4. DOI: 10.14722/ndss.2017.23160. URL: https://www.ndss-symposium.org/ndss2017/ndss-2017-programme/wireguard-next-generation-kernel-network-tunnel/ (visited on 09/28/2022) (cit. on p. 9).

[Fon22]   Juan Font. *Juanfont/Headscale*. Dec. 6, 2022. URL: https://github.com/juanfont/headscale (visited on 12/06/2022) (cit. on p. 10).

[Ope22]   OpenVPN. *Community Resources*. OpenVPN. 2022. URL: https://openvpn.net/community-resources/ (visited on 11/30/2022) (cit. on p. 9).

[Per18]   Trevor Perrin. "The Noise Protocol Framework". In: (July 1, 2018) (cit. on pp. 4, 10).

[Sli22]   Guus Sliepen. *Tinc Docs*. Nov. 30, 2022. URL: https://www.tinc-vpn.org/docs/ (visited on 11/30/2022) (cit. on p. 10).

[Tai]   Tailscale. *Tailscale*. Tailscale. URL: https://tailscale.com/kb/ (visited on 11/30/2022) (cit. on p. 10).