



Department of Mathematics and Computer Science
Coding Theory and Cryptology Group

Secure Sessions for Ad Hoc Multiparty Computation in MPyC

Master thesis

Emil Nikolov

Id nr: 0972305
`emil.e.nikolov@gmail.com`

Supervisor : Dr. ir. L.A.M. (Berry) Schoenmakers

May 9, 2023

Contents

Contents	iii
List of Figures	vii
1 Introduction	1
2 Related work	3
2.1 The Internet	3
2.1.1 Communications	3
2.1.2 Host Addressing	5
2.1.3 Security	7
2.2 Network overlays	8
2.2.1 TUN/TAP driver	8
2.2.2 Traditional VPNs	8
2.2.3 WireGuard	10
2.2.4 Mesh VPNs	11
2.2.5 Layer 7 overlays	11
3 Testing methodology	13
3.1 Measuring performance	13
3.2 Security	14
3.3 Usability	14
4 Internet Protocol based solution	15
4.1 Implementation details	15
4.2 Performance analysis	15
4.3 Security analysis	15
4.4 Usability analysis	15
5 WireGuard based solution	17
5.1 Implementation details	17
5.2 Performance analysis	17
5.3 Security analysis	17
5.4 Usability analysis	17
6 Tailscale based solution	19

CONTENTS

6.1	Implementation details	19
6.2	Performance analysis	19
6.3	Security analysis	19
6.3.1	Trust model	19
6.3.2	Identity	19
6.4	Usability analysis	19
7	Headscale based solution	21
7.1	Implementation details	21
7.2	Performance analysis	21
7.3	Security analysis	21
7.3.1	Trust model	21
7.3.2	Identity	21
7.4	Usability analysis	21

List of Abbreviations

E³ Extensible Evaluation Environment. 13

ACL Access Control List. 11

CGNAT Carrier-Grade NAT. 7

CIDR Classless Inter-Domain Routing. 4

DERP Designated Encrypted Relay for Packets. 6, 11

DNS Domain Name System. 5

ECDH Elliptic Curve Diffie-Hellman. 7

HTTP HyperText Transfer Protocol. 4

IP Internet Protocol. 3, 10

IPSec Internet Protocol Security. 7, 8

ISP Internet Service Provider. 3, 7

LAN Local Area Network. 3, 10

NAT Network Address Translation. 5

NAT-PMP NAT Port Mapping Protocol. 6

P2P Peer to Peer. 10, 11

PCP Port Control Protocol. 6

SSL Secure Sockets Layer. 7

STUN Session Traversal Utilities for NAT. 6

TCP Transmission Control Protocol. 4

TLS Transport Layer Security. 7

TURN Traversal Using Relays around NAT. 6

UDP User Datagram Protocol. 4

UPnP Universal Plug and Play. 6

URL Universal Resource Locator. 4

VPN Virtual Private Network. 7, 10, 11

WWW World Wide Web. 3

List of Figures

2.1	OSI model mapping of the Internet Protocol Suite	4
2.2	Two parties behind separate NATs	6
2.3	NAT traversal via STUN	7
2.4	OSI model mapping of various protocols	9

Chapter 1

Introduction

Chapter 2

Related work

This chapter explores prior work related to multi-party connectivity over The Internet and their suitability for joint MPyC computations. Section 2.1 provides an overview of the protocols of the modern Internet, the challenges for Peer-to-Peer communications and some of the approaches to overcome them. Section 2.2 discusses higher-level network overlays that use different combinations of the lower-level protocols from the previous section.

2.1 The Internet

The Internet is a global multi-tiered computer network of billions of host devices that communicate using the protocols of the Internet Protocol Suite (TCP/IP). Internet Service Providers (ISPs) are responsible for managing different sections of the infrastructure that connects the Local Area Networks (LANs) of various end-users including households and enterprises. The World Wide Web (WWW) or simply The Web is a collection of interconnected documents, e.g. HTML Web Pages, available on The Internet and is typically accessed via a user-agent software such as a **Web Browser**. The term “The Web” is sometimes used interchangeably with The Internet, but The Internet supports other services as well, e.g. file transfer (FTP), email (IMAP), instant messaging, remote access (SSH) and others. Conceptual frameworks like the OSI model are useful for understanding the objectives and functions of communication protocols. Figure 2.1 describes the responsibilities of the 7 layers of the OSI model and how they relate to the TCP/IP model used by the Internet Protocol Suite. The newer TCP/IP model only recognizes 4 layers as it merges the OSI Session (L5) and Presentation (L6) layers into the Application layer (L7), as well as the Physical layer (L1) into the Data link layer (L2). While the TCP/IP model is a more accurate representation of The Internet, the 7-layer numbers of the OSI model are still widely used in the literature.

2.1.1 Communications

The **Internet Protocol (IP)** is a Network (L3) protocol of the Internet Protocol Suite that is responsible for delivering datagrams between devices across the boundaries of their LANs by possibly routing traffic via multiple intermediate devices (routers). A datagram is a connectionless packet that is delivered on a best-effort basis. IP datagrams have a header that contains fields such as the **IP addresses** of its source and destination, and a payload that

show a diagram of the internet with multiple local networks

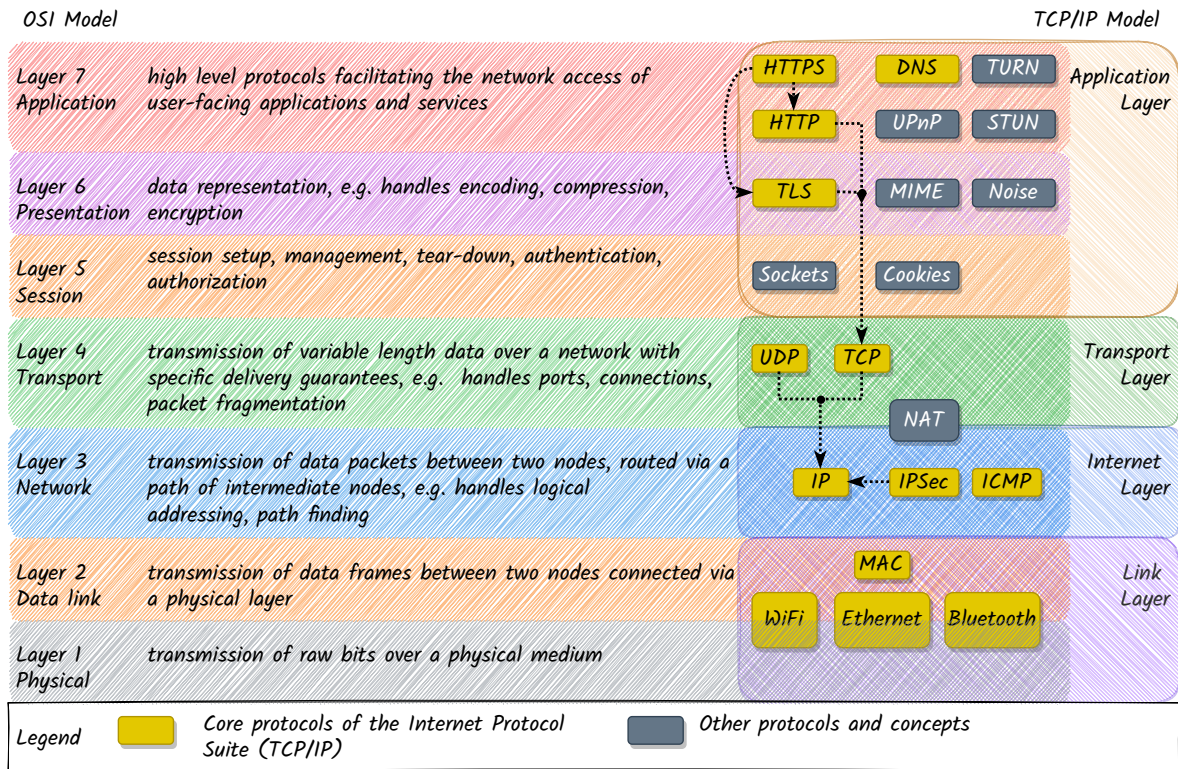


Figure 2.1: OSI model mapping of the Internet Protocol Suite

encapsulates the data from the protocols of the layers above. Routers are devices that are part of multiple networks and relay datagrams between them based on a routing table that maps IP address ranges to networks. *Classless Inter-Domain Routing (CIDR)* is a common notation for describing IP address ranges, e.g. 192.168.0.1/16, where the number after the slash describes the bit-length of the fixed prefix for a subnet.

User Datagram Protocol (UDP) and Transmission Control Protocol (TCP) are Transport layer (L4) protocols that employ 16-bit port numbers to enable multiple applications on the same host to establish their own communication channels while sharing an IP address. UDP offers faster communication, but only provides best-effort delivery, while TCP is a reliable transport protocol with stronger delivery guarantees at the expense of higher network latency. TCP maintains stateful connections that handle error detection and correction, packet ordering, flow control, acknowledgments and retransmissions in case packets are lost during transmission.

HyperText Transfer Protocol (HTTP) is an Application layer (L7) protocol used for transmitting data on The Web between web servers and clients like browsers, typically using a stateless request/response model. Similar to other L7 protocols, it uses *Universal Resource Locators (URLs)* for identifying resources using the format `scheme://host:port/path?query=value#fragment`, e.g. `http://www.example.com:80/path/to/file.html`. It is built on top of TCP and provides several features such as:

- Request Methods - used by the client to specify the action to perform on the resource behind the given URL, e.g. GET, POST, PUT, DELETE, etc.

- Headers - used to provide additional information about a request or response, e.g. Content-Type, Authorization, Cache-Control
- Status codes - used to indicate whether or not the result of a request, e.g. if it was successful (200), or if the resource is missing (404)
- Cookies - used to include stateful information about the user kept at the client-side
- Caching - used to specify that the result of a request can be cached for a certain time to avoid repeating the request's action.

The *Domain Name System (DNS)* operates at the Application Layer (L7) and allows the conversion of human-readable domains to IP addresses, e.g. `google.com` to `142.250.179.142`.

2.1.2 Host Addressing

The version of the Internet Protocol, that was originally deployed globally (IPv4), uses 32-bit numbers as IP addresses, allowing for around 4 billion unique addresses. Due to the popularity of The Internet, there are many more devices than available IPv4 addresses, which has caused challenges. IPv6 is a newer version of the protocol that uses a larger 128-bit address space which is sufficient for assigning 100 addresses for each atom on Earth. Its adoption has been slow, as according to Google as of 2023 only around 37% of the requests to their services use IPv6. Additionally, despite that IPv6 allows for all devices to be addressable on the Internet, for security reasons, most of them would use firewalls to block incoming remote traffic that is not associated with outgoing connections.

A widespread solution to the addressing problem is **Network Address Translation (NAT)**. It allows many devices without globally unique IP addresses to initiate connections to publicly addressable devices on the Internet via a limited number of gateways that must have globally unique IP addresses. A NAT gateway replaces the local source IP address of each outgoing IP datagram with its own public IP address before passing it on to the next link on the way to the destination while maintaining a mapping between the source and destination IPs in a translation table. The destination host can then address its responses back to the NAT gateway's public IP address, which in turn replaces its own IP from the incoming datagrams with the IP of the local device and forwards them to it. If the IP datagrams encapsulate TCP/UDP packets, the gateway additionally rewrites the source and destination ports, which means that NAT techniques can be placed somewhere between Layers 3 and 4 of the OSI model.

The effect of NAT on connectivity is similar to an IPv6 firewall as they both allow devices on a local network to initiate bidirectional communication to remote devices with public IP addresses, but connections cannot be natively initiated by the remote devices. As Figure 2.2 shows, it follows that when two devices are behind separate NATs, neither can contact the other first. **Client/Server** communication is less affected by this limitation because Servers are usually deployed to a public IP address that can be contacted by Clients with local IP addresses. **Peer-to-Peer** communication, however, is more challenging because the peers are often devices in separate residential networks behind different NATs. Several **NAT traversal** techniques try to solve this with different performance tradeoffs and success that varies depending on the NAT [**natBehaviorRFC**] and its behavior when mapping ports and IP addresses.

One approach based on the Client/Server model is to use a publicly addressable **relay** server

describe some of the nat behaviors, e.g. if the source IP address/port vary per destination are changed depending on the destination/-port mapping algorithms, if it




Figure 2.2: Two parties behind separate NATs

that is contacted by the NATed devices and then forwards the Peer-to-Peer traffic to the intended recipient. Compared to direct communication, relaying results in a higher network latency due to the longer path that each packet must travel. Maintaining a relay server requires some technical expertise and may be costly depending on the expected throughput. Despite the drawbacks, relaying works under most networking scenarios and is therefore often used as a fallback in case all other approaches fail to find a direct path. Protocols such as **Traversal Using Relays around NAT (TURN)** [[turnRFC](#)] and **Designated Encrypted Relay for Packets (DERP)** [[derpDocs](#)] can be used to securely implement relaying.

The NAT gateway in many residential networks is a Router device under the customer's control that has a statically or dynamically assigned public IP address. Most routers can be manually configured through their admin page to forward all traffic that arrives at a given port to a specific device on the local network. Remote applications can then initiate a connection to the local device if they know the IP address of the router and the forwarded port. The manual configuration, however, can be inconvenient and many users may be unaware of that setting because it is not necessary for the more straightforward Client/Server communications. Some routers also support programmatic configuration of port forwarding via a Layer 7 protocol like **Universal Plug and Play (UPnP)** or its successors **NAT Port Mapping Protocol (NAT-PMP)** and **Port Control Protocol (PCP)**. However, these protocols are not always supported and are often disabled by the local network administrators due to security concerns related to bugs in their implementation, vulnerable IOT devices on the local network or malicious programs being able to expose local devices to the internet.

An efficient NAT traversal approach that works with some types of NATs is to use **Session Traversal Utilities for NAT (STUN)** [[stunRFC](#)] in combination with UDP hole punching (Figure 2.3). STUN is a protocol operating at Layer 7 that allows a client application to detect the presence of NAT gateways on the network path to a public STUN server, and identify their types and the public IP address that they map to externally. An application sends UDP datagrams to the STUN server and it responds with the source IP address and port specified inside them. The application can compare its own endpoint with the source endpoint observed by the STUN server and if the values differ, it can be inferred that they were rewritten by a NAT. Additional STUN servers are contacted to determine if the NAT maps IPs and ports in a predictable fashion. UDP hole punching is a related technique that, depending on the NAT

types, can allow direct communication between two applications behind separate NATs. The applications must discover each other's externally mapped endpoints, perhaps via the STUN server. If the NATs use the same external port regardless of the remote destination, the two applications can simultaneously send UDP packets to each other's external endpoints. Their respective NATs will see the outgoing connection to the other peer - the "punched hole" - when the incoming traffic arrives from it and forward it correctly. NATs that map different ports per remote destination sometimes allocate port numbers predictably, which can be used by the peers to try to guess the port that will be opened by the opposing side's NAT.




Figure 2.3: NAT traversal via STUN

In mobile networks like 4G and 5G, the ISP often utilizes a **Carrier-Grade NAT (CGNAT)** as part of their infrastructure, while all devices under the user's control, including the router, only have local IP addresses. STUN techniques would fail to discover a direct path between two parties behind separate CGNATs or other unpredictable NAT algorithms. The only remaining possibility is to relay the traffic via a publicly reachable third-party host.

2.1.3 Security

Transport Layer Security (TLS) is a protocol that adds encryption on top of a reliable transport protocol such as TCP. It is a successor to Secure Sockets Layer (SSL). TLS does not strictly fit in any single OSI layer, but is usually placed somewhere between the Transport Layer (L4) and the Presentation Layer (L6). It is rather complex because it needs to support many possible use cases across the internet. The **Noise Protocol Framework** [Per18] is a more recent effort that applies the ideas of TLS in a simplified way by serving as a blueprint for designing use-case specific protocols for establishing secure communication channels based on Elliptic Curve Diffie-Hellman (ECDH) handshake patterns. It powers the end-to-end encryption in messaging applications such as WhatsApp and Signal, and Virtual Private Network (VPN) software such as WireGuard and Nebula.

HTTPS uses TLS to encrypt the HTTP traffic between a client and a server.

Internet Protocol Security (IPSec) is a protocol suite for encrypting the IP datagrams

only 65000 ports per IP address means that CGNATs that provide more than 65000 connections from client devices require more than one public IP address hairpinning - Hairpinning, also known as NAT loopback or NAT reflection, is a technique used by NAT devices to allow hosts on a private network to access a public server using its public IP address. Without hairpinning, the NAT device would not recognize the

between two hosts. It was originally developed as part of IPv6 but can also be used with IPv4. IPSec is similar in purpose to TLS but operates at the Network Layer (L3).

- **Internet Protocol Security (IPSec)**

- Layer 3 protocol suite part of the Internet Protocol Suite
- used inside VPN software
- has implementations in both user and kernel space as well as hardware implementations
- rewrites and encrypts the IP headers and payloads
- virtual routing table
- Initially was built into IPv6, separate from IPv4

2.2 Network overlays

Most Network Overlay solutions use a combination of the NAT traversal techniques mentioned previously. They can be placed in Layers 2, 3 or 7. Layer 2 overlays act as a virtual network switch, while Layer 3 overlays act as a virtual network router. Layer 7 overlays are implemented in user-space as libraries or applications that run on top of the network stack of the host operating system. Layer 2 and 3 overlays can either be implemented as kernel modules or as user-space applications that use a **TUN/TAP** driver to interface with the kernel.

Figure 2.4 shows an approximate OSI model mapping of several protocols and network overlay solutions from the point of view of the systems that use them and the arrows show dependency relations between them.

2.2.1 TUN/TAP driver

- **Layer 2 vs Layer 3 Networks**

- Layer 2 overlays bridge networks
 - * virtual network switch
 - * remote machines are on the same virtual LAN and can share the same IP address range
 - * allows broadcast/multicast
 - * TAP driver
- Layer 3 overlays route traffic between separate local networks
 - * virtual network router
 - * remote machines are on separate LANs
 - * simpler to configure
 - * TUN driver

2.2.2 Traditional VPNs

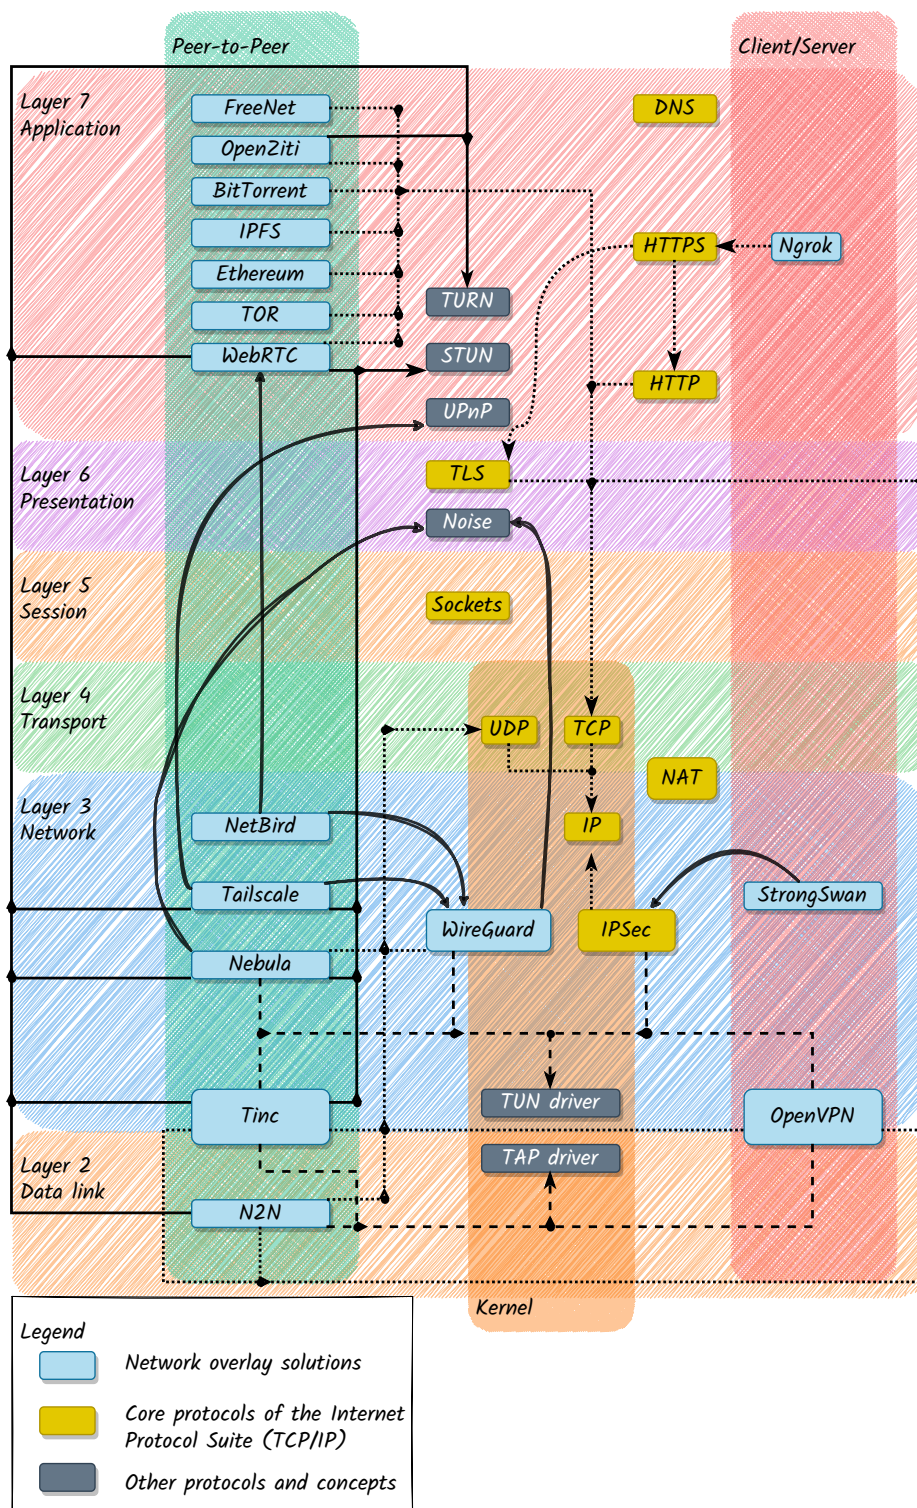


Figure 2.4: OSI model mapping of various protocols

VPNs are implemented as Layer 2 or 3 network overlays. They are commonly used for securely connecting machines from different LANs. They provide software emulation of a network interface controller via a TUN/TAP driver on the operating system level and allow other software to transparently use the functionality of the Internet Protocol (IP) suite without requiring extra changes. Traditional VPNs such as IPsec [ipSecDocs] and OpenVPN [Ope22] use a centralized service that all (encrypted) client communications must pass through. This introduces a single point of failure and a potential bottleneck that might negatively impact the performance of the multi-party computations due to their Peer to Peer (P2P) nature.

2.2.3 WireGuard

WireGuard [Don17] is a more recent protocol with a design informed by lessons learned from IPsec and OpenVPN and a key management approach inspired by SSH. It is a lower-level protocol that focuses on configuration simplicity while network topology, peer discovery and key distribution are left as a responsibility of higher-level systems that use it as a building block. Wireguard is implemented as a Layer 3 overlay over UDP tunnels. WireGuard has both user-space implementations that use a TUN driver and also has direct support built into the Linux Kernel since version 5.6 (May 2020). The kernel implementation allows for better performance because it does not need to copy packets between the kernel and user-space memory.

The snippets below show a minimal set of configuration options that need to be provided for two peers to be able to form secure tunnels with each other.

```
1 # peer1.conf
2 [Interface]
3 Address = 101.0.0.1/32
4 ListenPort = 53063
5 PrivateKey = ePTiXXhHjvAHdWUr8Bimk30n0gh3m241RAzsNOJZDW0=
6
7 [Peer]
8 PublicKey = BSn0ejd1Y3bKuD+Xpg0ZZe0f+Ies/oq10NZxw+S0mkc=
9 AllowedIPs = 101.0.0.2/32
10 Endpoint = peer1.example.com:38133
```

```
1 # peer2.conf
2 [Interface]
3 Address = 101.0.0.2/32
4 ListenPort = 38133
5 PrivateKey = sN/d6XUPEVPGSziVgCC0n0ivDK+qAoYC3nXnssQ5Rls=
6
7 [Peer]
8 PublicKey = e/TxvPmrgcc1G4cSH2bHv5J0PRHXKjYxTFoU8r+G93E=
9 AllowedIPs = 101.0.0.1/32
```

Each peer has a public/private key pair that is used for authentication and encryption based on the Noise Protocol Framework [Per18]. The **Address** field specifies the virtual IP address that the local network interface will use, while the **AllowedIPs** field specifies what virtual IP addresses are associated with a peer's public key. A peer's **Endpoint** field specifies the URL at which it can be reached. Only one of the peers must be configured with a reachable endpoint for the other one. In the above example once **peer1** initiates communication with

`peer2`, `peer2` will learn the current endpoint of `peer1` and will be able to communicate back with it.

2.2.4 Mesh VPNs

- Tinc
- N2N
- Tailscale
- Nebula
- ZeroTier

Mesh VPNs such as Tinc [Sli22], Tailscale [Tai] and Nebula [Def22] utilize NAT Traversal techniques to create direct P2P links between the clients for the data traffic. Authentication, authorization and traffic encryption are performed using certificates based on public key cryptography.

All three are open-source, except Tailscale's coordination service which handles peer discovery and identity management. Headscale [Fon22] is a community-driven open-source alternative for that component. Tinc is the oldest of the three but has a relatively small community. It is mainly developed by a single author and appears to be more academic than industry motivated. Nebula and Tailscale are both business driven. Tailscale was started by some high-profile ex-googlers and is the most end-user-focused of the three, providing a service that allows people to sign up using identity providers such as Google, Microsoft, GitHub and others. They also provide an Admin console that allows a user to easily add their personal devices to a network or share them with others. It also has support for automation tools like Terraform for creating authorization keys and managing an Access Control List (ACL) based firewall. Nebula was originally developed at the instant messaging company Slack to create overlay networks for their cross-region cloud infrastructure, but the authors later started a new company and are currently developing a user-centric platform similar to Tailscale's. Nebula is more customizable than Tailscale and since it is completely open-source it can be adapted to different use cases, but it is also more involved to set up. A certificate authority needs to be configured for issuing the identities of the participating hosts. Furthermore, publicly accessible coordination servers need to be deployed to facilitate the host discovery. Tailscale employs a distributed relay network of Designated Encrypted Relay for Packets (DERP) servers, while Nebula can be configured to route via one of the other peers in the VPN.

2.2.5 Layer 7 overlays

WebRTC is

- WebRTC
 - Uses STUN/TURN
 - d
- OpenZiti
 - uses relays
- ngrok
- TOR
- BitTorrent
- IPFS

- Ethereum
- Teleport
- Freenet

Chapter 3

Testing methodology

In the following chapters, we will design and implement several solutions for ad hoc MPC sessions based on a subset of the previously discussed related works:

- Internet protocol
- Wireguard
- Tailscale
- Headscale
- ? Headscale with DID identity?
- ? WebRTC?
- A custom solution that automates the WireGuard configuration by visiting a web page

Additionally, we will analyze and compare them in terms of performance, security and usability

3.1 Measuring performance

During the preparation phase of the project, we developed the Extensible Evaluation Environment (E^3) framework which simplifies and automates the process of deploying machines in different geographical regions, connecting them via an overlay network and executing multiparty computations between them, where each machine represents a different party.

To summarize, E^3 is a set of scripts that use several automation tools:

- Terraform - declarative provisioning
- NixOS - declarative Linux distribution
- Colmena - declarative deployment for NixOS
- PSSH - parallel execution of remote scripts over ssh
- DigitalOcean - a cloud provider

It allows us to quickly provision cloud virtual machines in multiple regions and reproducibly deploy all necessary software for running a multiparty computation over a chosen network overlay solution. The source code of E^3 can be found on [GitHub](#)

Each solution will be deployed using the E^3 framework and the performance will be quantitatively measured in terms of the time it takes to execute several MPyC demos. The selected

demos have different complexities in terms of communication rounds and message sizes which will allow us to observe their impact on the overall performance.

1. Secret Santa - high round complexity with small messages
2. Convolutional Neural Network (CNN) MNIST classifier - low round complexity with large messages

The demos will be configured at three different input size levels

- Low,
- Medium
- High

Furthermore, the demos will be executed in several networking scenarios:

1. 1-10 parties in the same geographic region
2. 1-10 parties evenly distributed across two nearby regions
3. 1-10 parties evenly distributed across two distant regions
4. 1-10 parties distributed across multiple distant regions

3.2 Security

We will analyze aspects such as

- key distribution
- trust model - are there any trusted third parties and what would be the consequences if they are corrupted or breached
- traffic encryption
- identity strength

3.3 Usability

For each solution, we will describe the steps that the parties need to perform to execute a joint multiparty computation. Those steps will be analyzed in terms of:

- Complexity - how much technical expertise is expected from the parties to be able to execute the steps
- Initial effort - how much effort is each party expected to put in preparing for their first joint computation
- Repeated effort - after the initial setup, how much effort is required to perform another computation
 - with the same set of parties
 - with another set of parties
- Finalization effort - how much effort is required to finalize the MPC session once it is complete and clean up any left-over artifacts or resources so that the machine of each party is in its original state

Chapter 4

Internet Protocol based solution

This solution focuses on directly using the internet protocol without involving an overlay network. Our goal is to analyze the implications of using only the functionalities that MPyC directly supports to serve as the reference for our later experiments.

4.1 Implementation details

We will manually set up the multiparty computations via the public IP addresses of the machines and DNS.

4.2 Performance analysis

4.3 Security analysis

4.4 Usability analysis

Chapter 5

WireGuard based solution

This solution creates an overlay network by manually configuring WireGuard on each machine.

5.1 Implementation details

5.2 Performance analysis

5.3 Security analysis

5.4 Usability analysis

Chapter 6

Tailscale based solution

Tailscale is a VPN solution that configures a mesh of direct Wireguard tunnels between the peers.

6.1 Implementation details

6.2 Performance analysis

6.3 Security analysis

6.3.1 Trust model

There is a centralized service that deals with the key distribution, which needs to be trusted to provide the correct public keys for the correct parties

6.3.2 Identity

Identity is based on third party identity providers such as Microsoft and GitHub

- Magic DNS
-

6.4 Usability analysis

With tailscale each party needs to

- register a Tailscale account
- Download and install tailscale on the machine they want to run a multiparty computation
- Run tailscale on their machine and logs into their account in order to link it to their own Tailnet
- Share their Tailscale machine with the Tailnets of each of the other parties
- Download the demo they want to run
- Form the flags for running the chosen demo

- add -P \$HOST:\$PORT for each party using their Tailscale hostname/virtual IP
- Run the demo

Chapter 7

Headscale based solution

This solution is similar to the Tailscale one, but it uses Headscale - a self-hosted open-source alternative to the closed-source Tailscale coordination service.

7.1 Implementation details

7.2 Performance analysis

7.3 Security analysis

7.3.1 Trust model

There still is a centralized service like in the Tailscale solution, but here it is self-deployed.

7.3.2 Identity

7.4 Usability analysis

Bibliography

- [Def22] Defined. *Nebula: Open Source Overlay Networking / Nebula Docs*. 2022. URL: <https://docs.defined.net/docs/> (visited on 12/01/2022) (cit. on p. 11).
- [Don17] Jason A. Donenfeld. “WireGuard: Next Generation Kernel Network Tunnel”. In: *Proceedings 2017 Network and Distributed System Security Symposium*. Network and Distributed System Security Symposium. San Diego, CA: Internet Society, 2017. ISBN: 978-1-891562-46-4. DOI: [10.14722/ndss.2017.23160](https://doi.org/10.14722/ndss.2017.23160). URL: <https://www.ndss-symposium.org/ndss2017/ndss-2017-programme/wireguard-next-generation-kernel-network-tunnel/> (visited on 09/28/2022) (cit. on p. 10).
- [Fon22] Juan Font. *Juanfont/Headscale*. Dec. 6, 2022. URL: <https://github.com/juanfont/headscale> (visited on 12/06/2022) (cit. on p. 11).
- [Ope22] OpenVPN. *Community Resources*. OpenVPN. 2022. URL: <https://openvpn.net/community-resources/> (visited on 11/30/2022) (cit. on p. 10).
- [Per18] Trevor Perrin. “The Noise Protocol Framework”. In: (July 1, 2018) (cit. on pp. 7, 10).
- [Sli22] Guus Sliepen. *Tinc Docs*. Nov. 30, 2022. URL: <https://www.tinc-vpn.org/docs/> (visited on 11/30/2022) (cit. on p. 11).
- [Tai] Tailscale. *Tailscale*. Tailscale. URL: <https://tailscale.com/kb/> (visited on 11/30/2022) (cit. on p. 11).